

Capston Project - Marketing campaign

February 4, 2026

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
# Python
import warnings

# Suppress all warnings
warnings.filterwarnings("ignore")
```

```
[2]: df= pd.read_csv("bank-full.csv", sep=";")
df.head()
```

```
[2]:    age          job marital education default balance housing loan \
0   58 management married  tertiary      no     2143    yes    no
1   44 technician single secondary      no      29    yes    no
2   33 entrepreneur married secondary      no       2    yes   yes
3   47 blue-collar married unknown      no    1506    yes    no
4   33        unknown single unknown      no       1    no    no

      contact day month duration campaign pdays previous poutcome y
0  unknown    5   may      261         1     -1        0  unknown  no
1  unknown    5   may      151         1     -1        0  unknown  no
2  unknown    5   may       76         1     -1        0  unknown  no
3  unknown    5   may       92         1     -1        0  unknown  no
4  unknown    5   may      198         1     -1        0  unknown  no
```

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
 #   Column      Non-Null Count Dtype  
--- 
 0   age         45211 non-null  int64  
 1   job          45211 non-null  object  
 2   marital     45211 non-null  object  
 3   education   45211 non-null  object  

```

```

4   default    45211 non-null  object
5   balance    45211 non-null  int64
6   housing    45211 non-null  object
7   loan       45211 non-null  object
8   contact    45211 non-null  object
9   day        45211 non-null  int64
10  month      45211 non-null  object
11  duration   45211 non-null  int64
12  campaign   45211 non-null  int64
13  pdays      45211 non-null  int64
14  previous   45211 non-null  int64
15  poutcome   45211 non-null  object
16  y          45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB

```

[]:

0.1 Univariate analysis for categorical and object variables

[4]: # Univariate analysis of categorical and object variables

```


def plot_object(dataframe, column_name):
    """
    Plots a bar chart showing category frequencies with both frequency (inside bar)
    and proportion (above bar) labels.
    Parameters:- dataframe: pandas DataFrame- column_name: str, name of the
    categorical column to visualize
    """
    # Count frequencies and proportions
    value_counts = dataframe[column_name].value_counts()
    proportions = value_counts / len(dataframe)
    # Set plot style
    sns.set(style="whitegrid")
    plt.figure(figsize=(10, 6))
    # Bar plot
    palette1=sns.color_palette(palette='tab20')
    ax = sns.barplot(x=value_counts.index, y=value_counts.values, palette=palette1)
    # Annotate bars
    for i, (count, prop) in enumerate(zip(value_counts.values, proportions.
    values)):
        # Frequency inside bar
        ax.text(i, count * 0.5, f'{count}', ha='center', va='center', fontsize=11,
        color='black', fontweight='bold')
        # Proportion above bar
        ax.text(i, count + max(value_counts.values) * 0.02, f'{prop:.2%}', ha='center', fontsize=10, color='black')

```

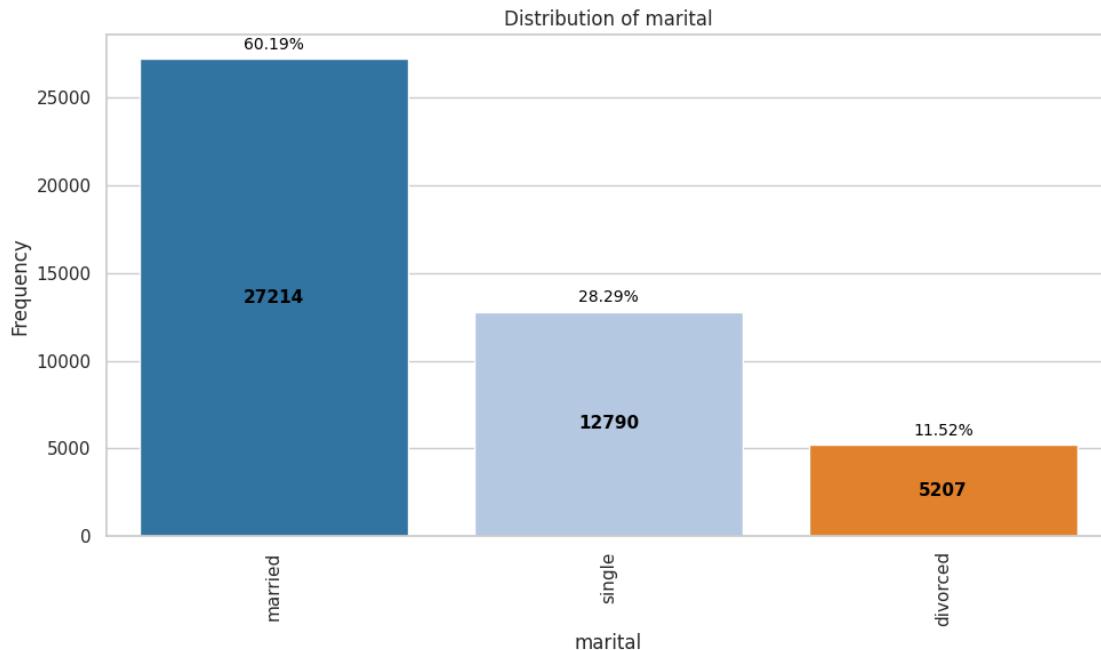
```
plt.title(f'Distribution of {column_name}')
plt.xlabel(column_name)
plt.xticks(rotation=90)
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```

[]:

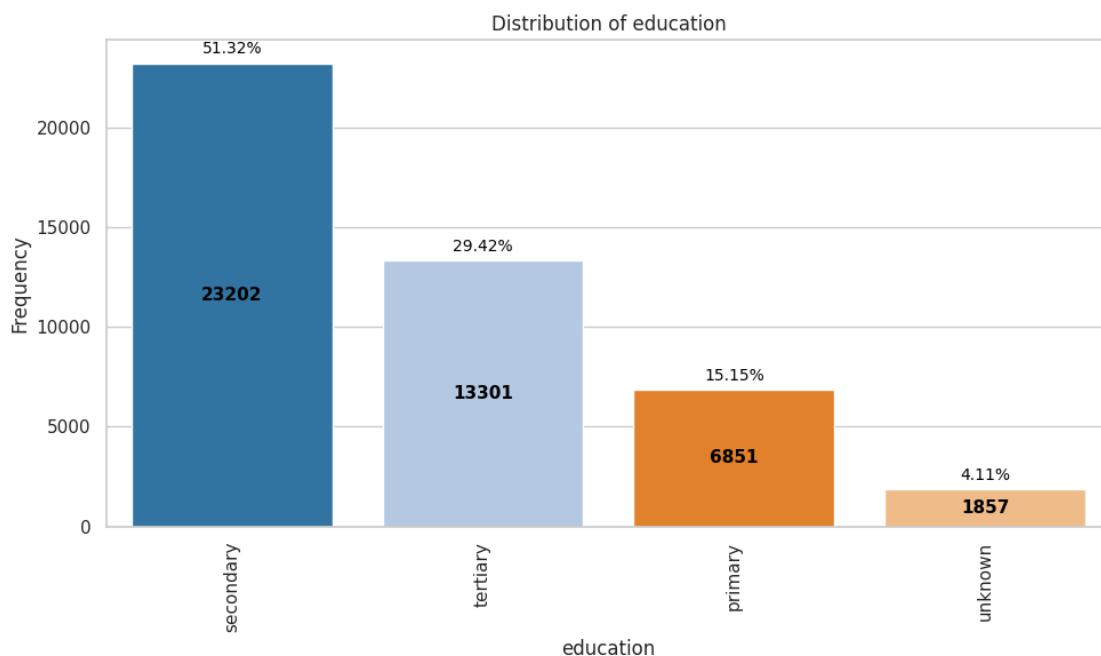
[5]: plot_object(df, "job")



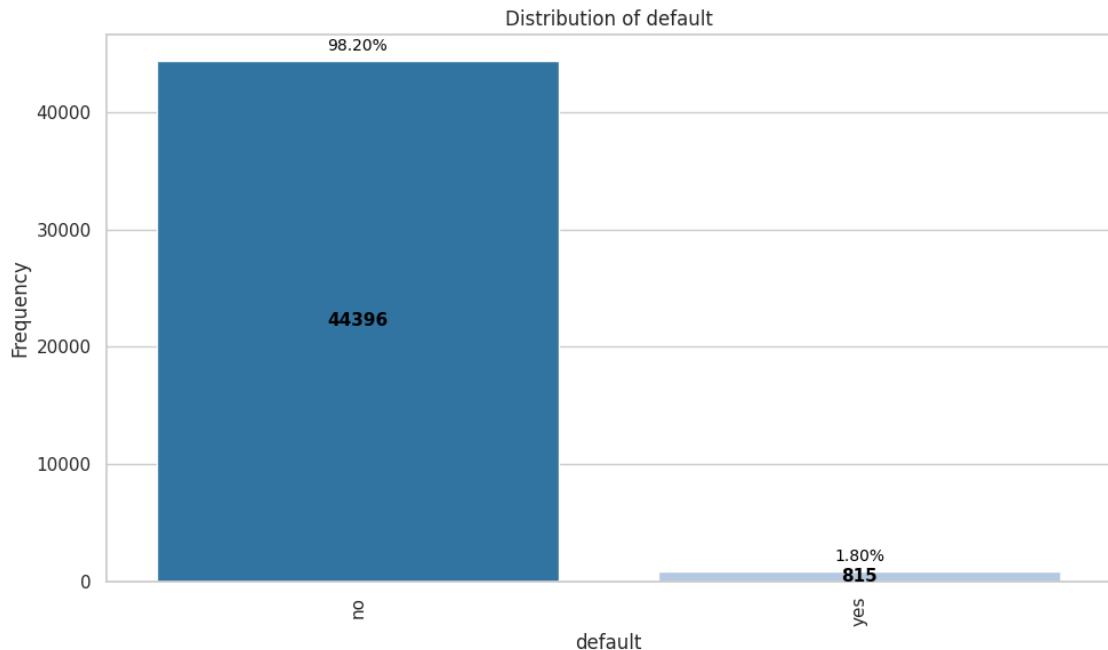
[6]: plot_object(df, 'marital')




```
[7]: plot_object(df, 'education')
```



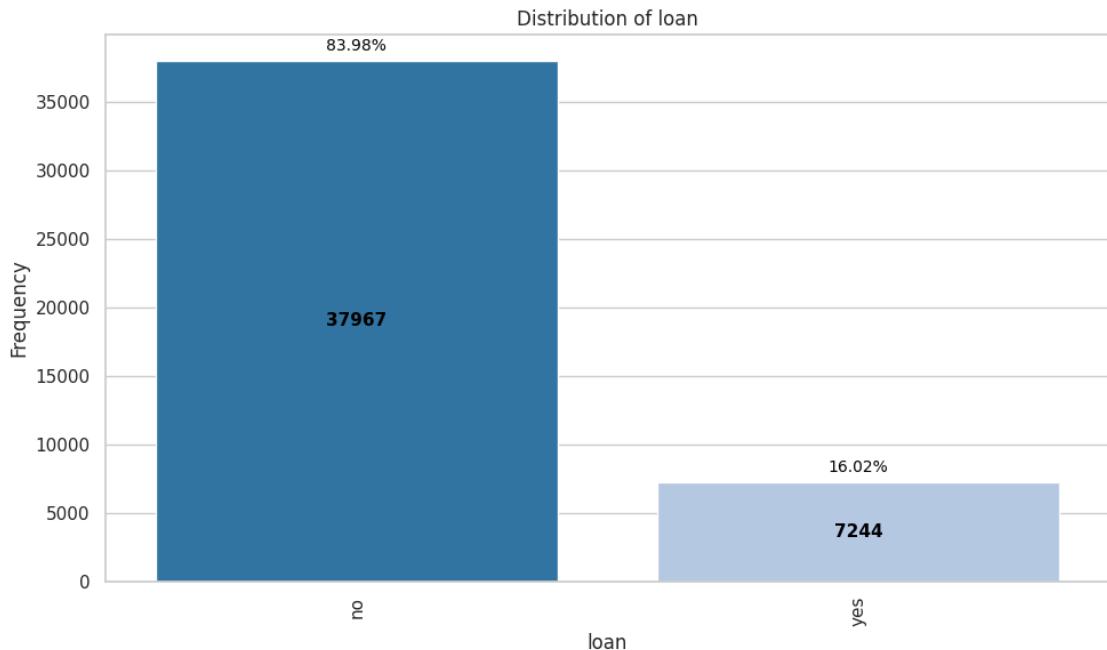
```
[8]: plot_object(df, 'default')
```



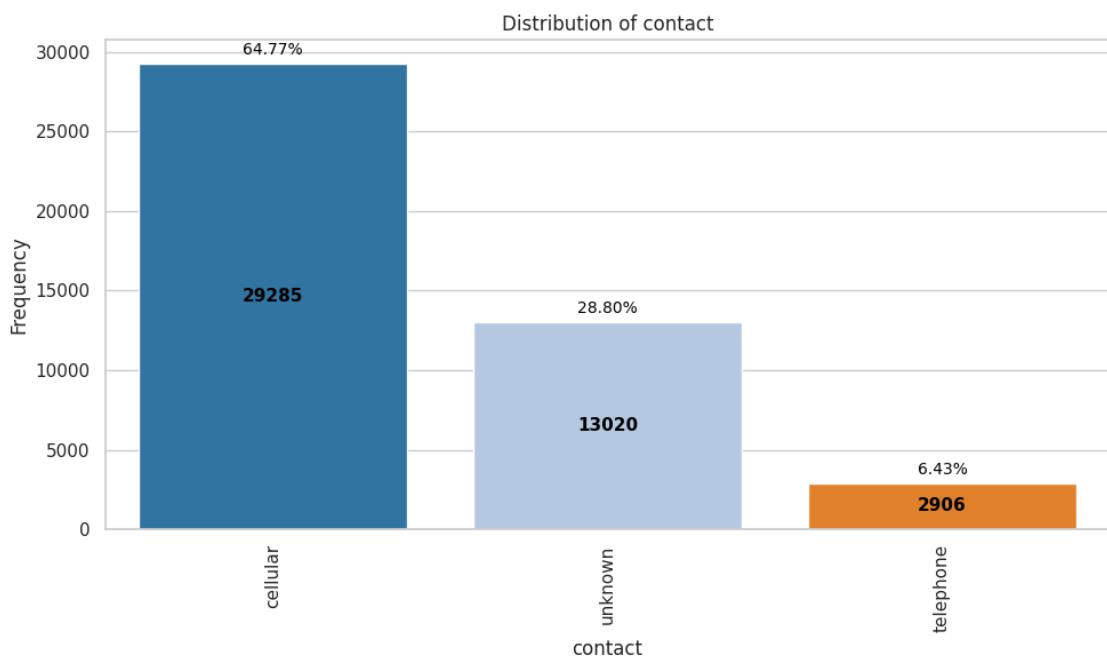
```
[9]: plot_object(df, 'housing')
```



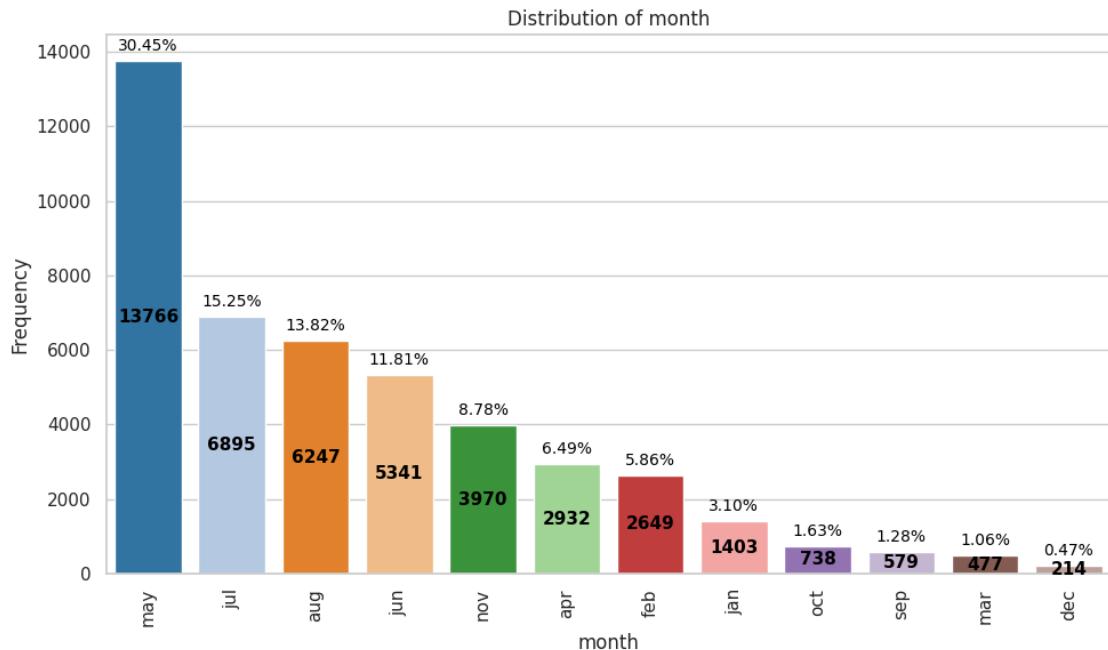
```
[10]: plot_object(df, 'loan')
```



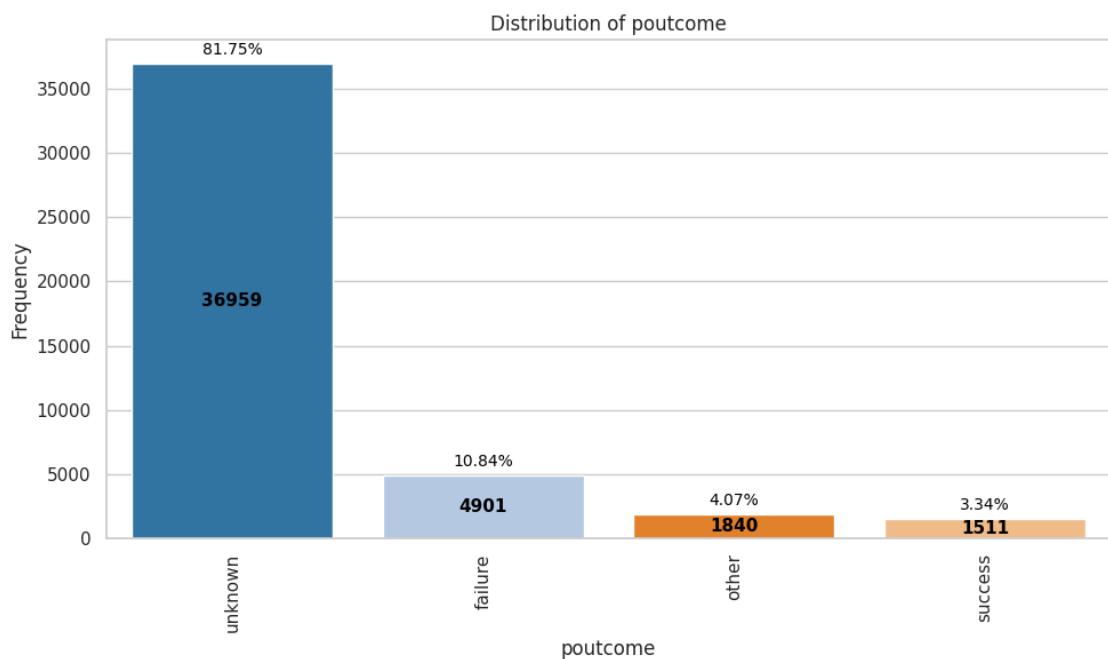
```
[11]: plot_object(df, 'contact')
```




```
[12]: plot_object(df, 'month')
```



```
[13]: plot_object(df, 'poutcome')
```



```
[14]: plot_object(df, 'y')
```




0.2 Univariate Analysis for numerical variables

```
[15]: # univariate analysis of continuous variables
def cont_plot(df, var):
    #var="Age"
    # Set plot style
    sns.set(style="whitegrid")
    # Create a figure with two subplots: histogram and box plot
    plt.figure(figsize=(12, 6))
    # Histogram
    # Box plot
    plt.subplot(1, 2, 1)
    sns.histplot(df[var], bins=20, kde=True, color='red')
    plt.title(var+' Distribution- Histogram')
    plt.xlabel(var)
    plt.ylabel('Frequency')
    plt.subplot(1, 2, 2)
    sns.boxplot(x=df[var], color='yellow')
    plt.title(var+' Distribution- Box Plot')
```


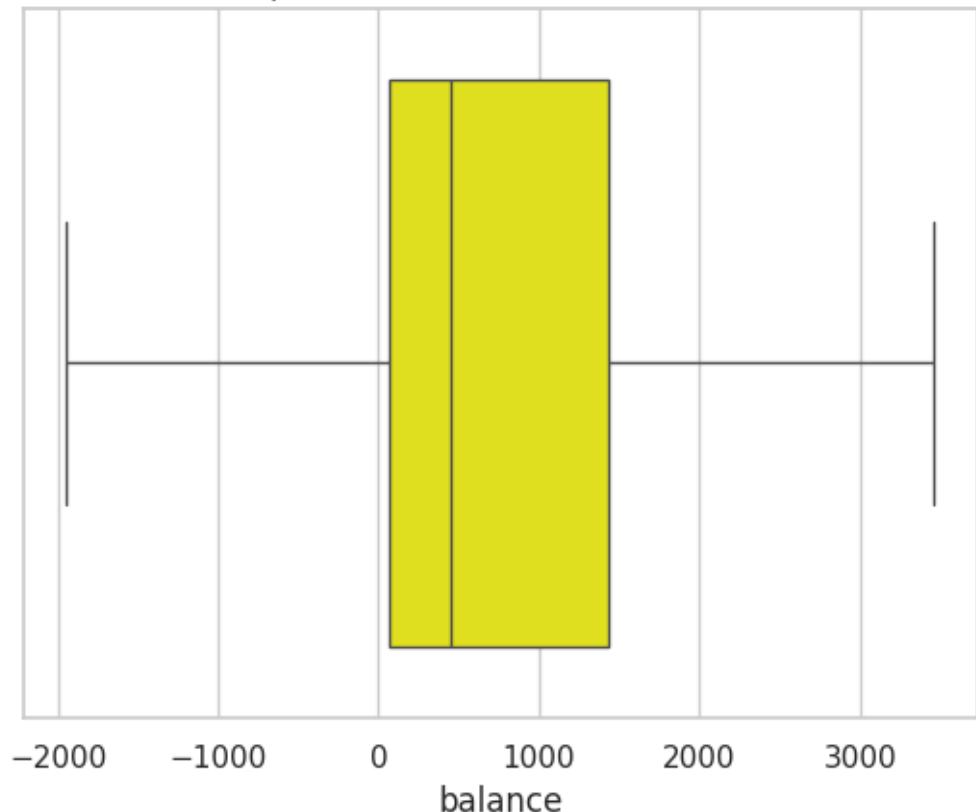
```
[16]: def NoOutlier(df,var):
    sns.boxplot(x=df[var], showfliers=False,color="yellow")
    plt.title("Boxplot of "+ var +" (without outliers)")
    plt.show()
```

```
[17]: cont_plot(df, 'age' )
```

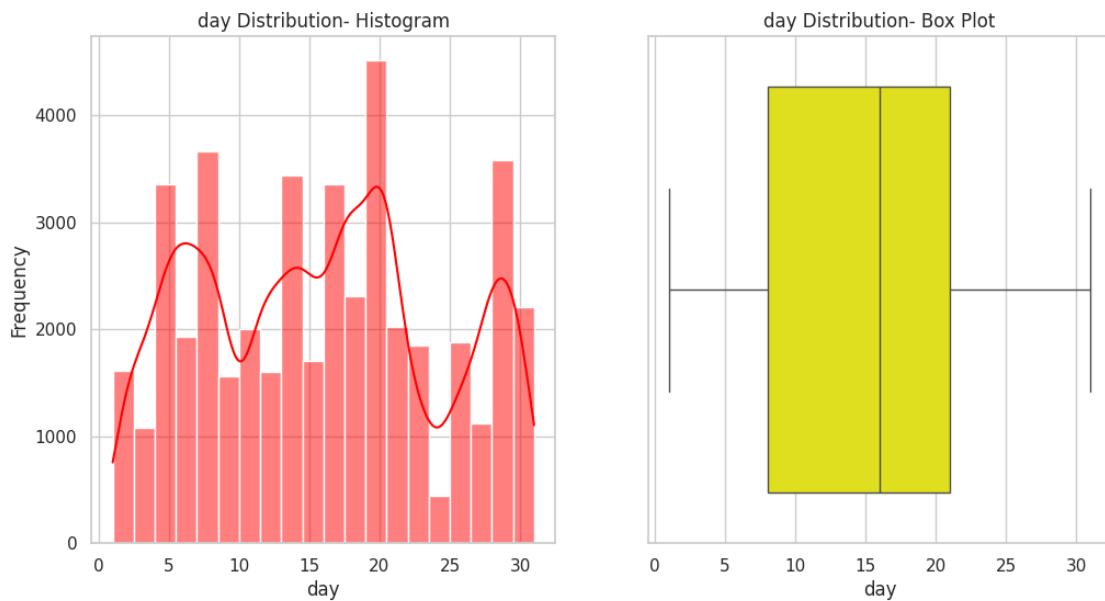


```
[18]: # Boxplot without outliers  
NoOutlier(df, "balance")  
cont_plot(df, 'balance' )
```

Boxplot of balance (without outliers)

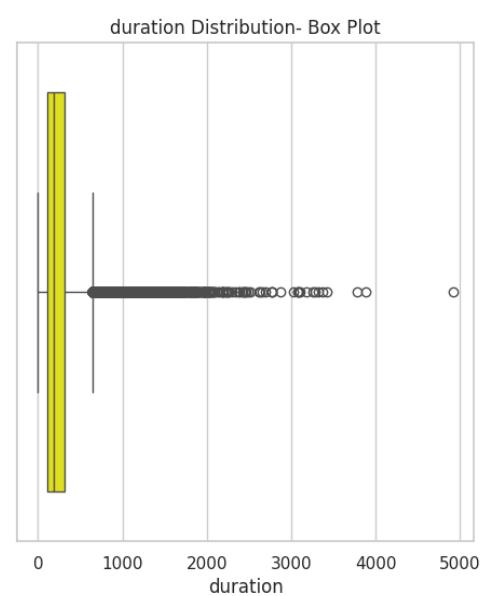
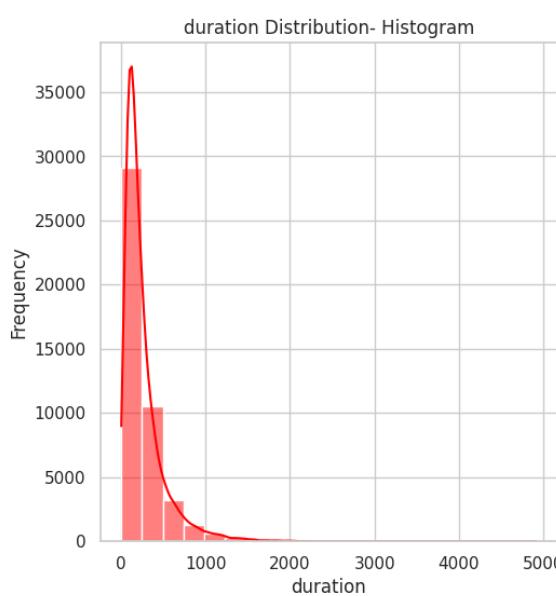
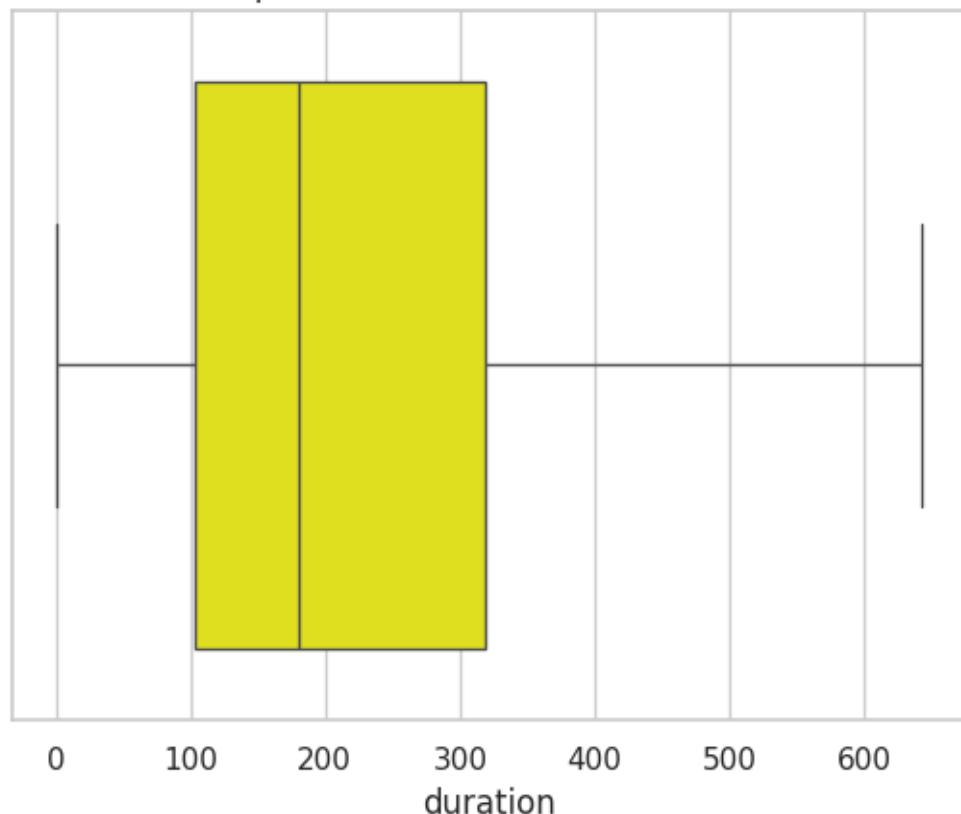


```
[19]: cont_plot(df, 'day' )
```




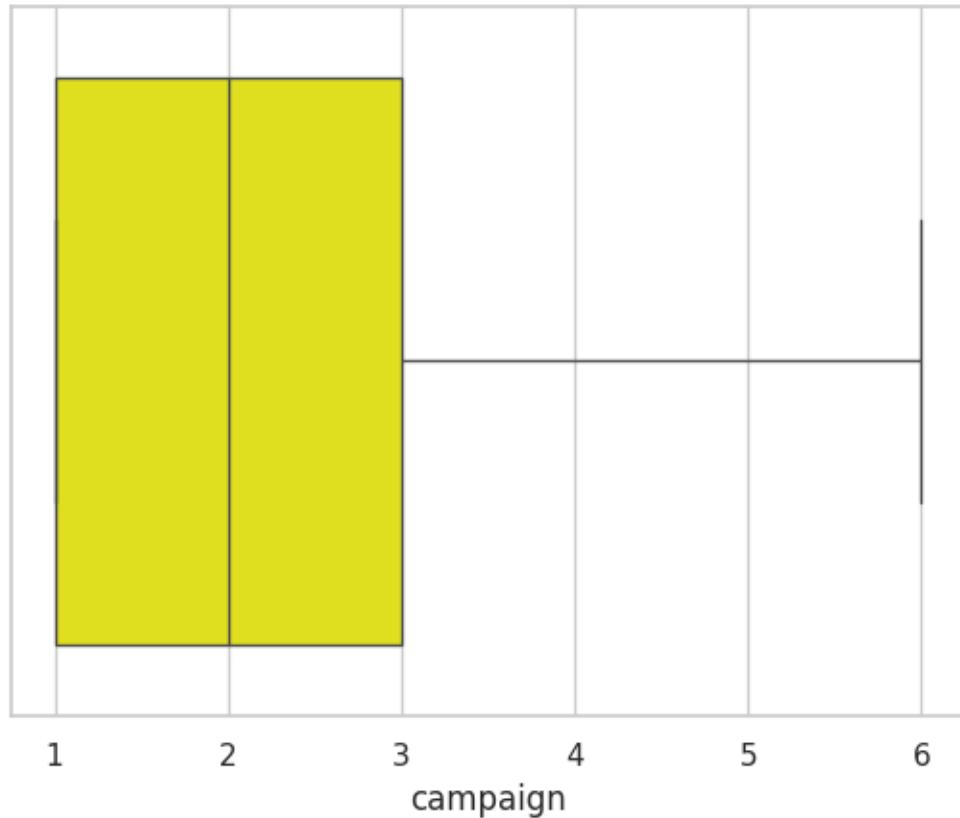
```
[20]: NoOutlier(df, "duration")
cont_plot(df, 'duration' )
```

Boxplot of duration (without outliers)




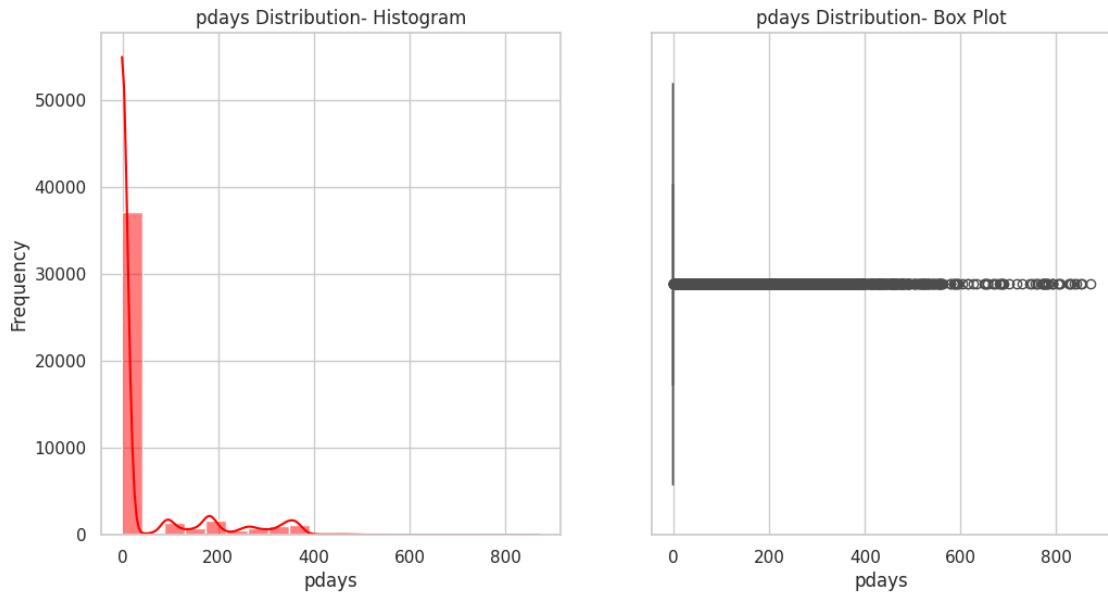
```
[21]: NoOutlier(df , "campaign")
cont_plot(df , 'campaign' )
```

Boxplot of campaign (without outliers)





```
[22]: NoOutlier(df,"pdays")
cont_plot(df,'pdays' )
```





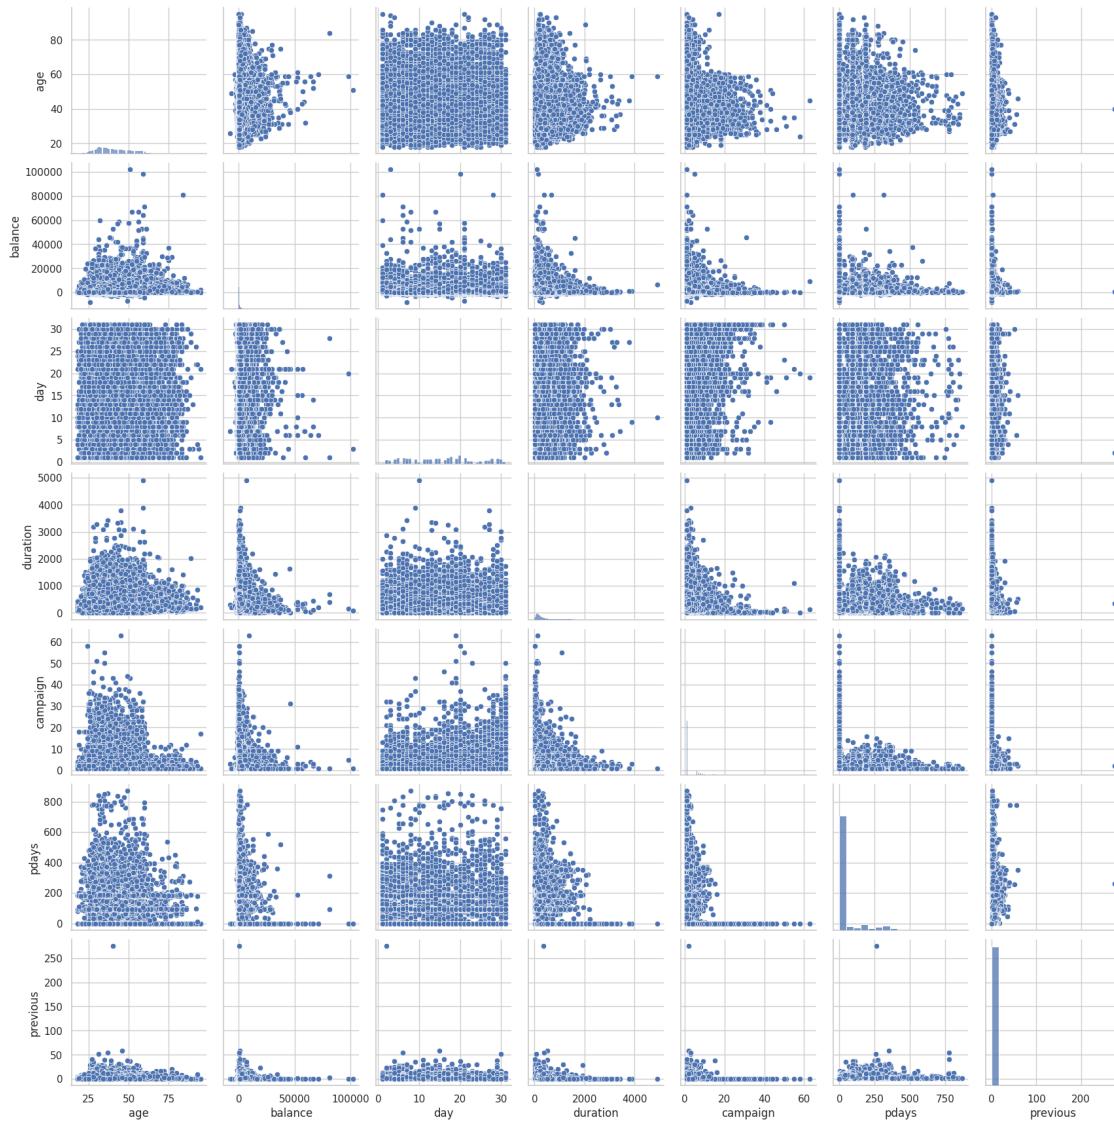
```
[23]: NoOutlier(df, "previous")
cont_plot(df, 'previous')
```





```
[24]: import seaborn as sns
import matplotlib.pyplot as plt


# Pairwise scatter plots for numerical variables
sns.pairplot(df.select_dtypes(include=['number']))
plt.show()
```



```
[25]: import seaborn as sns
import matplotlib.pyplot as plt

# Compute correlation matrix
corr_matrix = df.select_dtypes(include=['number']).corr()

# Plot heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0)
plt.title("Pairwise Correlation of Numerical Variables")
plt.show()
```



```
[26]: int_columns = df.select_dtypes('int64').columns.tolist()
print(int_columns)
```

```
['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous']
```

```
[ ]:
```

```
[27]: object_columns = df.select_dtypes(include=['object']).columns.tolist()
```

```
print(object_columns)
```

```
['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
'month', 'poutcome', 'y']
```


```
[28]: def cat_cont_plot(df, xvar, yvar):
    """
    Creates histograms of a continuous variable across categories.

    Parameters:
    - df: pandas DataFrame
    - xvar: str, categorical column name (e.g. 'Target')
    - yvar: str, continuous column name (e.g. 'Age')
    """
    sns.boxplot(x=xvar, y=yvar, data=df, color="pink")
    plt.xticks(rotation=90)


    # Create a larger facet grid of histograms by category
    #g = sns.FacetGrid(df, col=xvar, col_wrap=4, height=6, aspect=0.5) # ↴ Increased size and layout
    #g.map(sns.histplot, yvar, bins=20, color="green")

    # Add titles and labels
    #g.set_axis_labels(yvar, "Frequency")
    #g.set_titles(col_template="{col_name}")
    #plt.tight_layout()
    plt.show()
```

```
[29]: cat_cont_plot(df, "y", "duration")
```



```
[30]: cat_cont_plot(df, "y", "campaign")
```



```
[31]: def cat_cat_plot(dataframe, column_name, hue_column):
    """
    Plots a grouped bar chart showing category frequencies split by hue,
    with both frequency (inside bar) and proportion (above bar) labels.

    Parameters:
    - dataframe: pandas DataFrame
    - column_name: str, name of the categorical column to visualize (x-axis)
    - hue_column: str, name of the second categorical variable to group by (hue)
    """
    # Count combinations of column and hue
    counts_df = dataframe.groupby([column_name, hue_column]).size().
    ↪reset_index(name='count')
    total_counts = dataframe[column_name].value_counts()

    # Set plot style
    sns.set(style="whitegrid")
    plt.figure(figsize=(14, 8)) # Larger frame

    # Create bar plot
```


```

ax = sns.barplot(x=column_name, y='count', hue=hue_column, data=counts_df, palette="muted")
# Annotate bars
for container in ax.containers:
    for bar in container:
        height = bar.get_height()
        x = bar.get_x() + bar.get_width() / 2
        category = bar.get_label()
        base_x = int(round(x)) # used for proportion lookup
        if height > 0:
            ax.text(x, height * 0.5, f'{int(height)}', ha='center', va='center',
                     fontsize=10, color='black', fontweight='bold')
            ax.text(x, height + max(counts_df['count']) * 0.02, f'{height / sum(container.datavalues):.1%}',
                     ha='center', fontsize=10, color='navy', fontweight='bold')


# Beautify plot
plt.title(f'{column_name} Distribution by {hue_column}', fontsize=16)
plt.xlabel(column_name, fontsize=23)
plt.ylabel('Frequency', fontsize=19)
plt.xticks(rotation=45)
plt.legend(title=hue_column)
plt.tight_layout()
plt.show()

```

[32]: cat_cat_plot(df, 'job', "y")



```
[33]: cat_cat_plot(df,'education','y')
```



```
[ ]:
```

```
[34]: #import numpy as np

# Indicator: whether client was contacted before
df['previous_contact'] = (df['pdays'] != -1).astype(int)

# Replace -1 with NaN so pdays is only meaningful when contact exists
df['pdays'] = df['pdays'].replace(-1, np.nan)
```

1 Research Question 1: Which customer and campaign features best predict term deposit subscription?

```
[35]: rq1_vars=["age", "job", "education", "marital", "balance", "default", "housing", "loan", "contact", "duration", "poutcome", "previous", "campaign", "contact", "previous_contact", 'pdays', "y"]
```

```
[36]: dfrq1=df[rq1_vars]
```

```
[37]: dfrq1.head()
```

```
[37]:   age      job education marital balance default housing loan \
0    58  management    tertiary  married     2143      no     yes   no
1    44  technician  secondary   single       29      no     yes   no
2    33 entrepreneur  secondary  married        2      no     yes  yes
3    47 blue-collar    unknown  married     1506      no     yes   no
4    33      unknown    unknown   single        1      no      no   no

      contact duration poutcome previous campaign contact previous_contact \
0  unknown       261  unknown         0          1  unknown                  0
1  unknown       151  unknown         0          1  unknown                  0
2  unknown        76  unknown         0          1  unknown                  0
3  unknown       92  unknown         0          1  unknown                  0
4  unknown      198  unknown         0          1  unknown                  0

      pdays     y
0    NaN  no
1    NaN  no
2    NaN  no
3    NaN  no
4    NaN  no
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

1.1 Preprocessing

```
[38]: dfrq1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              45211 non-null   int64  
 1   job              45211 non-null   object  
 2   education        45211 non-null   object  
 3   marital          45211 non-null   object  
 4   balance          45211 non-null   int64  
 5   default          45211 non-null   object  
 6   housing          45211 non-null   object  
 7   loan              45211 non-null   object  
 8   contact          45211 non-null   object  
 9   duration          45211 non-null   int64  
 10  poutcome         45211 non-null   object  
 11  previous         45211 non-null   int64  
 12  campaign         45211 non-null   int64  
 13  contact          45211 non-null   object  
 14  previous_contact 45211 non-null   int64  
 15  pdays             8257 non-null   float64 
 16  y                 45211 non-null   object  
dtypes: float64(1), int64(6), object(10)
memory usage: 5.9+ MB
```

```
[ ]:
```

```
[ ]:
```

```
[39]: dfrq1.shape
```

```
[39]: (45211, 17)
```

1.1.1 Encoding all object variables into dummy variable.

```
[40]: cat_cols = dfrq1.select_dtypes(include='object').columns

dfrq1_enc = pd.get_dummies(dfrq1, columns=cat_cols, drop_first=True)

print(dfrq1_enc.shape)
print(dfrq1_enc.head())
```

```
(45211, 38)
age  balance  duration  previous  campaign  previous_contact  pdays  \

```

```

0    58      2143      261      0      1      0      NaN
1    44       29      151      0      1      0      NaN
2    33        2       76      0      1      0      NaN
3    47     1506      92      0      1      0      NaN
4    33        1      198      0      1      0      NaN

    job_blue-collar  job_entrepreneur  job_housemaid ... contact_telephone \
0                      0                  0                  0   ...
1                      0                  0                  0   ...
2                      0                  1                  0   ...
3                     1                  0                  0   ...
4                      0                  0                  0   ...

    contact_unknown  poutcome_other  poutcome_success  poutcome_unknown \
0                     1                  0                  0      1
1                     1                  0                  0      1
2                     1                  0                  0      1
3                     1                  0                  0      1
4                     1                  0                  0      1

    contact_telephone  contact_unknown  contact_telephone  contact_unknown \
0                      0                  1                  0      1
1                      0                  1                  0      1
2                      0                  1                  0      1
3                      0                  1                  0      1
4                      0                  1                  0      1

y_yes
0    0
1    0
2    0
3    0
4    0

```

[5 rows x 38 columns]

[41]: `from sklearn.model_selection import train_test_split`

```
X = dfrq1_enc.drop(columns=['y_yes'])
y = dfrq1_enc['y_yes']
```

[42]: `#Train-test split (stratified)`
`from sklearn.model_selection import train_test_split`

```
X_train, X_test, y_train, y_test = train_test_split( X, y,
    test_size=0.3,
    stratify=y,
```

```
    random_state=42  
)
```

Given the presence of substantial outliers in several numerical variables, RobustScaler was used to scale the features, as it relies on the median and interquartile range and is less sensitive to extreme values.

[]:

[]:

[]:

Since pdays is undefined for clients never previously contacted, a binary indicator was introduced to represent prior contact status. Remaining missing values in pdays were imputed using the median to avoid introducing artificial extremes.

The variable pdays uses the value -1 to indicate clients who were never previously contacted. As this value is not a valid numerical quantity, it was replaced with missing values, and a binary indicator variable (previous_contact) was created to capture prior contact status. Since logistic regression does not handle missing values natively, a median imputation strategy was applied within a modeling pipeline prior to robust scaling and model fitting. RobustScaler was used due to the presence of substantial outliers in several numerical variables.

```
[43]: from sklearn.pipeline import Pipeline  
from sklearn.impute import SimpleImputer  
from sklearn.preprocessing import RobustScaler  
from sklearn.linear_model import LogisticRegression  
  
log_reg_pipeline = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='median')),      # handles NaNs (pdays)  
    ('scaler', RobustScaler()),                          # robust to outliers  
    ('model', LogisticRegression(  
        max_iter=1000,  
        class_weight='balanced',  
        random_state=42  
    ))  
])
```

```
[44]: log_reg_pipeline.fit(X_train, y_train)
```

```
[44]: Pipeline(steps=[('imputer', SimpleImputer(strategy='median')),  
                     ('scaler', RobustScaler()),  
                     ('model',  
                      LogisticRegression(class_weight='balanced', max_iter=1000,  
                                         random_state=42))])
```

Model evaluation

```
[45]: from sklearn.metrics import confusion_matrix, classification_report

y_pred = log_reg_pipeline.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[10002  1975]
 [ 341  1246]]
```

```
[46]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.97	0.84	0.90	11977
1	0.39	0.79	0.52	1587
accuracy			0.83	13564
macro avg	0.68	0.81	0.71	13564
weighted avg	0.90	0.83	0.85	13564

```
[47]: #ROC-AUC
from sklearn.metrics import roc_auc_score

y_prob = log_reg_pipeline.predict_proba(X_test)[:, 1]
roc_auc = roc_auc_score(y_test, y_prob)

print("ROC-AUC:", roc_auc)
```

ROC-AUC: 0.89182053883049

Interpret coefficients

```
[48]: import pandas as pd
import numpy as np

coef = log_reg_pipeline.named_steps['model'].coef_[0]

coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': coef,
    'Odds_Ratio': np.exp(coef)
}).sort_values(by='Odds_Ratio', ascending=False)

coef_df
```

[48]:

	Feature	Coefficient	Odds_Ratio
31	poutcome_success	2.385626	10.865859
2	duration	1.193241	3.297752
14	job_student	0.568677	1.765929
19	education_tertiary	0.518236	1.679063
11	job_retired	0.424058	1.528151
20	education_unknown	0.307635	1.360205
30	poutcome_other	0.260160	1.297138
18	education_secondary	0.258276	1.294695
22	marital_single	0.189902	1.209131
5	previous_contact	0.177928	1.194740
1	balance	0.040230	1.041050
3	previous	0.030356	1.030821
35	contact_telephone	0.012216	1.012291
33	contact_telephone	0.012216	1.012291
28	contact_telephone	0.012216	1.012291
26	contact_telephone	0.012216	1.012291
0	age	0.009416	1.009460
6	pdays	0.000241	1.000241
32	poutcome_unknown	-0.101671	0.903327
21	marital_married	-0.178142	0.836823
10	job_management	-0.203363	0.815982
4	campaign	-0.236401	0.789464
23	default_yes	-0.252914	0.776535
16	job_unemployed	-0.272010	0.761847
15	job_technician	-0.322706	0.724186
36	contact_unknown	-0.332182	0.717357
29	contact_unknown	-0.332182	0.717357
34	contact_unknown	-0.332182	0.717357
27	contact_unknown	-0.332182	0.717357
17	job_unknown	-0.345373	0.707957
13	job_services	-0.448088	0.638848
8	job_entrepreneur	-0.529650	0.588811
9	job_housemaid	-0.533286	0.586674
7	job_blue-collar	-0.536965	0.584520
12	job_self-employed	-0.538566	0.583585
25	loan_yes	-0.687911	0.502625
24	housing_yes	-0.849760	0.427518

[]:

Interpretation rule:

Odds Ratio > 1 → increases probability of subscription

Odds Ratio < 1 → decreases probability

```
[49]: #Extract coefficients and odds ratios from the pipeline
import pandas as pd
import numpy as np


# Extract coefficients
coef = log_reg_pipeline.named_steps['model'].coef_[0]

coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': coef,
    'Odds_Ratio': np.exp(coef)
})

coef_df = coef_df.sort_values(by='Odds_Ratio', ascending=False)
```

```
[50]: #Visualise coefficients (direction + strength)
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 10))
plt.barh(coef_df['Feature'], coef_df['Coefficient'])
plt.axvline(0)
plt.xlabel('Logistic Regression Coefficient')
plt.title('Logistic Regression Coefficients')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```




```
[51]: #Visualise odds ratios
plt.figure(figsize=(8, 10))
plt.barh(coef_df['Feature'], coef_df['Odds_Ratio'])
plt.axvline(1)
plt.xlabel('Odds Ratio')
```

```

plt.title('Odds Ratios from Logistic Regression')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```



2 Random forest

```
[52]: X = dfrq1_enc.drop(columns=['y_yes'])
y = dfrq1_enc['y_yes']
```

```
[53]: from sklearn.model_selection import train_test_split

X_train_rf, X_test_rf, y_train_rf, y_test_rf = train_test_split(
    X,
    y,
    test_size=0.3,
    stratify=y,
    random_state=42
)
```

Fit Random Forest (no scaling needed)

Random Forest does NOT need scaling and handles outliers naturally.

```
[54]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(
    n_estimators=300,
    max_depth=None,
    min_samples_leaf=5,
    class_weight='balanced',
    random_state=42,
    n_jobs=-1
)

rf.fit(X_train_rf, y_train_rf)
```

```
[54]: RandomForestClassifier(class_weight='balanced', min_samples_leaf=5,
                             n_estimators=300, n_jobs=-1, random_state=42)
```

Model evaluation

```
[55]: from sklearn.metrics import confusion_matrix, classification_report,
      roc_auc_score

y_pred = rf.predict(X_test)
y_prob = rf.predict_proba(X_test)[:, 1]

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_prob))
```

Confusion Matrix:

```
[[10494  1483]
 [ 384 1203]]
```

```

Classification Report:
      precision    recall  f1-score   support

          0       0.96     0.88      0.92     11977
          1       0.45     0.76      0.56      1587

   accuracy                           0.86     13564
macro avg       0.71     0.82      0.74     13564
weighted avg    0.90     0.86      0.88     13564

```

ROC-AUC: 0.9014899329996021

Feature importance

```
[56]: import pandas as pd

feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': rf.feature_importances_
}).sort_values(by='Importance', ascending=False)

feature_importance.head(15)
```

```
[56]:      Feature  Importance
2        duration  0.468353
1        balance  0.078820
0           age  0.073920
31  poutcome_success  0.051599
6            pdays  0.042453
24        housing_yes  0.035750
4        campaign  0.035113
27  contact_unknown  0.020728
34  contact_unknown  0.020588
29  contact_unknown  0.020351
36  contact_unknown  0.017481
3        previous  0.015889
25        loan_yes  0.010693
21  marital_married  0.010467
5  previous_contact  0.009960
```

```
[57]: #Visualise top feature importances
import matplotlib.pyplot as plt


top_features = feature_importance.head(15)

plt.figure(figsize=(8, 8))
plt.barh(top_features['Feature'], top_features['Importance'])
```

```

plt.xlabel('Feature Importance')
plt.title('Top 15 Feature Importances (Random Forest)')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```



A Random Forest classifier was trained to identify key customer and campaign features associated with term deposit subscription. As a tree-based model, Random Forest does not require feature scaling and is robust to outliers. Class imbalance was addressed using balanced class weights. Model performance was evaluated using recall, F1-score, and ROC–AUC. Feature importance scores were extracted to identify the most influential predictors.

2.1 SHAP

To enhance interpretability of the Random Forest model, SHAP (SHapley Additive exPlanations) values were computed. SHAP provides a model-agnostic explanation framework that quantifies both the magnitude and direction of each feature's contribution to the predicted outcome. The SHAP summary plot highlights the most influential features and reveals whether higher or lower values of each feature increase the likelihood of subscription.

```
[58]: # Take a random sample of the test set
X_test_shap = X_test.sample(n=1000, random_state=42)
```

```
[59]: import shap

explainer = shap.TreeExplainer(rf)

shap_values = explainer.shap_values(X_test_shap)
```


```
[64]: import shap

# Use the unified SHAP API
explainer = shap.Explainer(rf, X_train)

# Compute SHAP values on the sample
shap_values = explainer(X_test_shap)

# Summary plot (this WILL work)
shap.summary_plot(
    shap_values.values,
    X_test_shap,
    plot_type="dot",
    show=True
)
```

100%|=====| 1998/2000 [02:38<00:00]



[]:

[69]:

[68]: # -----
5. Compute SHAP INTERACTION values (slow but needed once)

```
shap_interaction_values = explainer.shap_interaction_values(X_test_shap)
```

```

# For binary classification: take positive class ("yes")
shap_inter_pos = shap_interaction_values[:, :, 1]

# -----
# 6. Quantify MAIN vs INTERACTION effects
# -----
# ---- MAIN EFFECTS (positive class only) ----
shap_main = np.mean(
    np.abs(shap_values.values[:, :, 1]), # class "yes"
    axis=0
)
# Mean absolute MAIN effects per feature
#shap_main = np.mean(np.abs(shap_values.values), axis=0)

# Mean absolute INTERACTION effects
shap_inter = np.mean(np.abs(shap_inter_pos), axis=0)

# Remove self-interactions (diagonal)
np.fill_diagonal(shap_inter, 0)

# Aggregate interaction strength per feature
interaction_strength = shap_inter.sum(axis=1)

# -----
# 7. Comparison table (proof of weak interactions)
# -----
comparison = pd.DataFrame({
    "feature": X_test_shap.columns,
    "main_effect": shap_main,
    "interaction_effect": interaction_strength,
    "interaction_ratio": interaction_strength / shap_main
}).sort_values("interaction_ratio", ascending=False)

comparison

```

FEATURE_DEPENDENCE::independent does not support interactions!

	feature	main_effect	interaction_effect	interaction_ratio
0	age	0.016193	0.0	0.0
19	education_tertiary	0.011398	0.0	0.0
21	marital_married	0.010235	0.0	0.0
22	marital_single	0.006799	0.0	0.0
23	default_yes	0.000343	0.0	0.0
24	housing_yes	0.034003	0.0	0.0
25	loan_yes	0.010002	0.0	0.0
26	contact_telephone	0.000414	0.0	0.0

27	contact_unknown	0.014227	0.0	0.0
28	contact_telephone	0.000433	0.0	0.0
29	contact_unknown	0.013769	0.0	0.0
30	poutcome_other	0.001024	0.0	0.0
31	poutcome_success	0.017082	0.0	0.0
32	poutcome_unknown	0.007596	0.0	0.0
33	contact_telephone	0.000420	0.0	0.0
34	contact_unknown	0.013588	0.0	0.0
35	contact_telephone	0.000501	0.0	0.0
20	education_unknown	0.000665	0.0	0.0
18	education_secondary	0.002449	0.0	0.0
1	balance	0.020787	0.0	0.0
17	job_unknown	0.000015	0.0	0.0
2	duration	0.134506	0.0	0.0
3	previous	0.008536	0.0	0.0
4	campaign	0.017766	0.0	0.0
5	previous_contact	0.008049	0.0	0.0
6	pdays	0.013353	0.0	0.0
7	job_blue-collar	0.005185	0.0	0.0
8	job_entrepreneur	0.000691	0.0	0.0
9	job_housemaid	0.000449	0.0	0.0
10	job_management	0.003651	0.0	0.0
11	job_retired	0.001034	0.0	0.0
12	job_self-employed	0.000531	0.0	0.0
13	job_services	0.001159	0.0	0.0
14	job_student	0.001493	0.0	0.0
15	job_technician	0.002351	0.0	0.0
16	job_unemployed	0.000379	0.0	0.0
36	contact_unknown	0.012287	0.0	0.0

[]:

Although Random Forest models are capable of capturing nonlinear interactions, SHAP interaction analysis shows that interaction effects are negligible for all predictors. The model's predictions are therefore driven primarily by additive contributions of individual features rather than by complex feature interactions.

Methods / Explainability section

SHAP interaction values were computed on a representative subset of the test data to assess whether the Random Forest model relied on nonlinear feature interactions. Interaction effects were quantified and compared to main feature effects using the ratio of interaction strength to main effect magnitude.

Results / Interpretation section

Across all predictors, interaction effects were effectively zero relative to main effects. This indicates that the Random Forest model predominantly relies on additive feature contributions, with no evidence of strong pairwise interactions influencing subscription predictions.

This finding suggests that the predictive structure of the data is largely linear-additive, despite the use of a non-linear model.

[]: