

Full code-Capston Project - Marketing campaign

February 18, 2026

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import RobustScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report,
    roc_auc_score, accuracy_score
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
import shap
# Python
import warnings

# Suppress all warnings
warnings.filterwarnings("ignore")
```

```
[2]: df= pd.read_csv("bank-full.csv", sep=";")
df.head()
```

```
[2]:   age          job marital education default  balance housing loan \
0    58 management married  tertiary     no      2143     yes    no
1    44 technician single secondary    no       29     yes    no
2    33 entrepreneur married secondary    no        2     yes   yes
3    47 blue-collar married unknown     no     1506     yes    no
4    33         unknown single unknown     no         1     no    no

  contact day month duration campaign pdays previous poutcome    y
0 unknown    5   may      261           1      -1        0  unknown  no
1 unknown    5   may      151           1      -1        0  unknown  no
```

```

2    unknown      5   may        76       1     -1          0   unknown   no
3    unknown      5   may        92       1     -1          0   unknown   no
4    unknown      5   may       198       1     -1          0   unknown   no

```

[3]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
 #   Column      Non-Null Count Dtype  
---  --  
 0   age         45211 non-null  int64  
 1   job          45211 non-null  object 
 2   marital      45211 non-null  object 
 3   education    45211 non-null  object 
 4   default      45211 non-null  object 
 5   balance      45211 non-null  int64  
 6   housing      45211 non-null  object 
 7   loan          45211 non-null  object 
 8   contact      45211 non-null  object 
 9   day           45211 non-null  int64  
 10  month         45211 non-null  object 
 11  duration     45211 non-null  int64  
 12  campaign     45211 non-null  int64  
 13  pdays         45211 non-null  int64  
 14  previous     45211 non-null  int64  
 15  poutcome     45211 non-null  object 
 16  y             45211 non-null  object 
dtypes: int64(7), object(10)
memory usage: 5.9+ MB

```

[]:

0.1 Univariate analysis for categorical and object variables

[4]: # Univariate analysis of categorical and object variables

```

def plot_object(dataframe, column_name):
    """
    Plots a bar chart showing category frequencies with both frequency (inside bar)
    and proportion (above bar) labels.
    Parameters:- dataframe: pandas DataFrame- column_name: str, name of the
    categorical column to visualize
    """
    # Count frequencies and proportions
    value_counts = dataframe[column_name].value_counts()
    proportions = value_counts / len(dataframe)

```


```

# Set plot style
sns.set(style="whitegrid")
plt.figure(figsize=(10, 6))
# Bar plot
palette1=sns.color_palette(palette='tab20')
ax = sns.barplot(x=value_counts.index, y=value_counts.values, palette=palette1)
# Annotate bars
for i, (count, prop) in enumerate(zip(value_counts.values, proportions.
values)):
    # Frequency inside bar
    ax.text(i, count * 0.5, f'{count}', ha='center', va='center', fontsize=11, color='black', fontweight='bold')
    # Proportion above bar
    ax.text(i, count + max(value_counts.values) * 0.02, f'{prop:.2%}', ha='center', fontsize=10, color='black')
plt.title(f'Distribution of {column_name}')
plt.xlabel(column_name)
plt.xticks(rotation=90)
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()


```

[]:


[5]: plot_object(df, "job")




```
[6]: plot_object(df, 'marital')
```




```
[7]: plot_object(df, 'education')
```




```
[8]: plot_object(df, 'default')
```




```
[9]: plot_object(df, 'housing')
```




```
[10]: plot_object(df, 'loan')
```




```
[11]: plot_object(df, 'contact')
```




```
[12]: plot_object(df, 'month')
```



```
[13]: plot_object(df, 'poutcome')
```



```
[14]: plot_object(df, 'y')
```




0.2 Univariate Analysis for numerical variables

```
[15]: # univariate analysis of continuous variables
def cont_plot(df, var):
    #var="Age"
    # Set plot style
    sns.set(style="whitegrid")
    # Create a figure with two subplots: histogram and box plot
    plt.figure(figsize=(12, 6))
    # Histogram
    # Box plot
    plt.subplot(1, 2, 1)
    sns.histplot(df[var], bins=20, kde=True, color='red')
    plt.title(var+' Distribution- Histogram')
    plt.xlabel(var)
    plt.ylabel('Frequency')
    plt.subplot(1, 2, 2)
    sns.boxplot(x=df[var], color='yellow')
    plt.title(var+' Distribution- Box Plot')
```

```
[16]: def NoOutlier(df,var):
    sns.boxplot(x=df[var], showfliers=False,color="yellow")
    plt.title("Boxplot of "+ var +" (without outliers)")
```

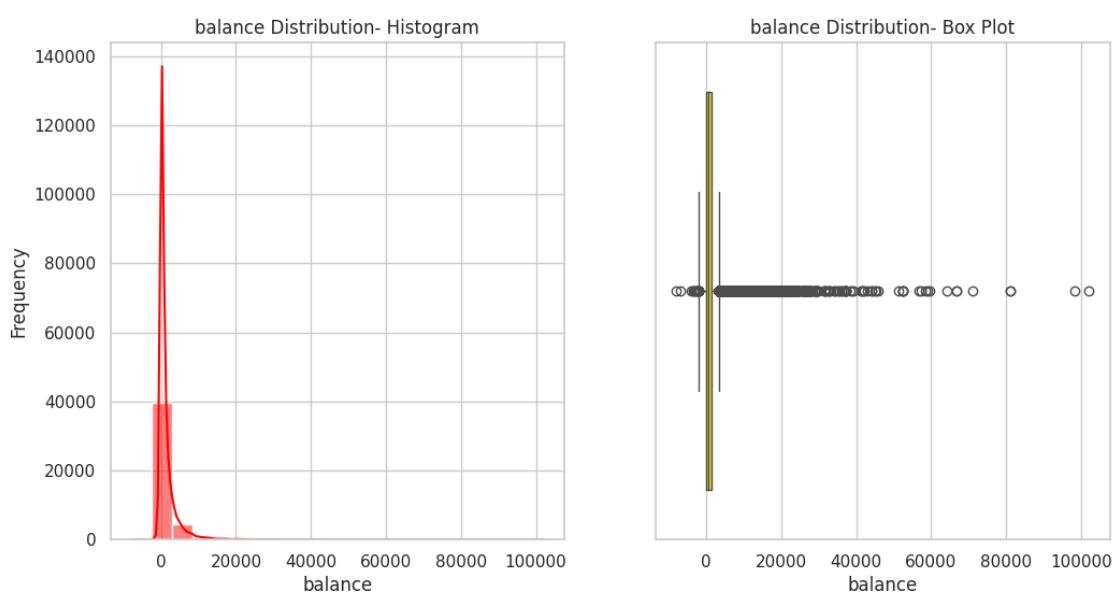
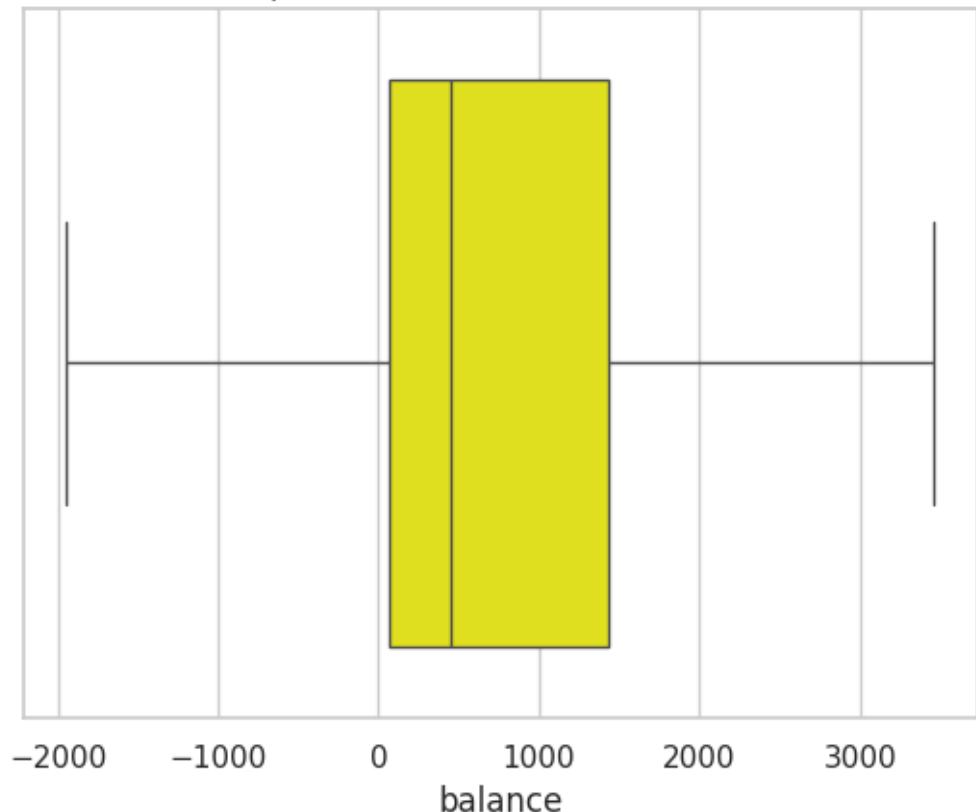
```
plt.show()
```

```
[17]: cont_plot(df, 'age' )
```

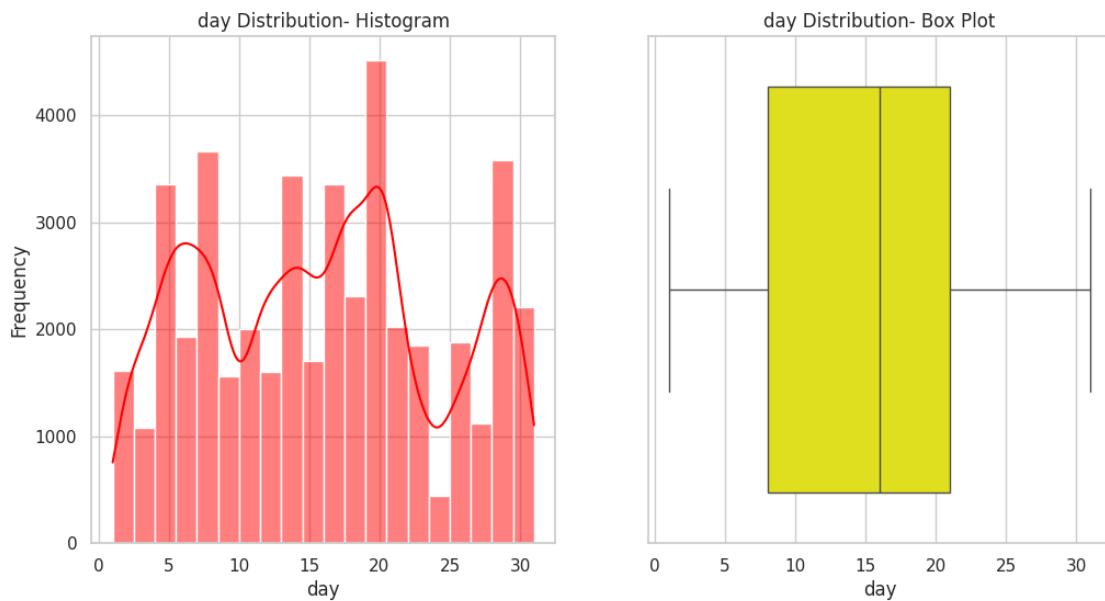


```
[18]: # Boxplot without outliers  
NoOutlier(df, "balance")  
cont_plot(df, 'balance' )
```

Boxplot of balance (without outliers)

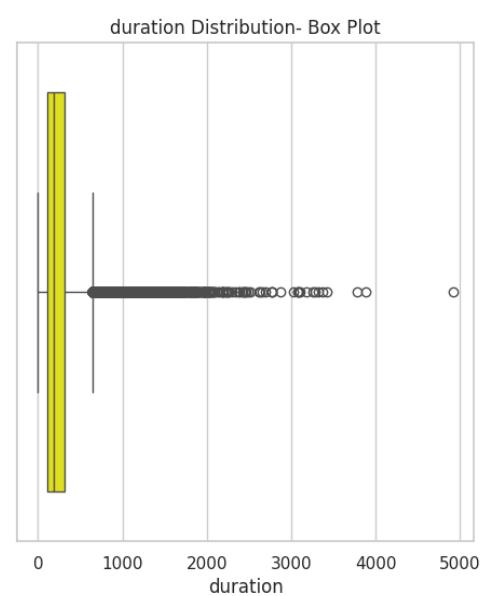
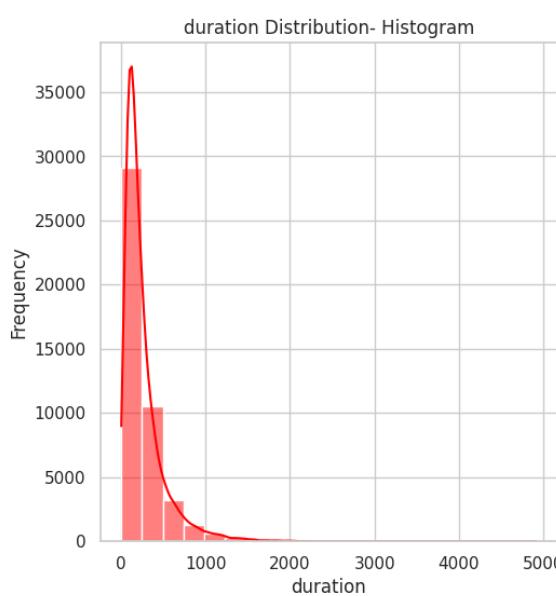
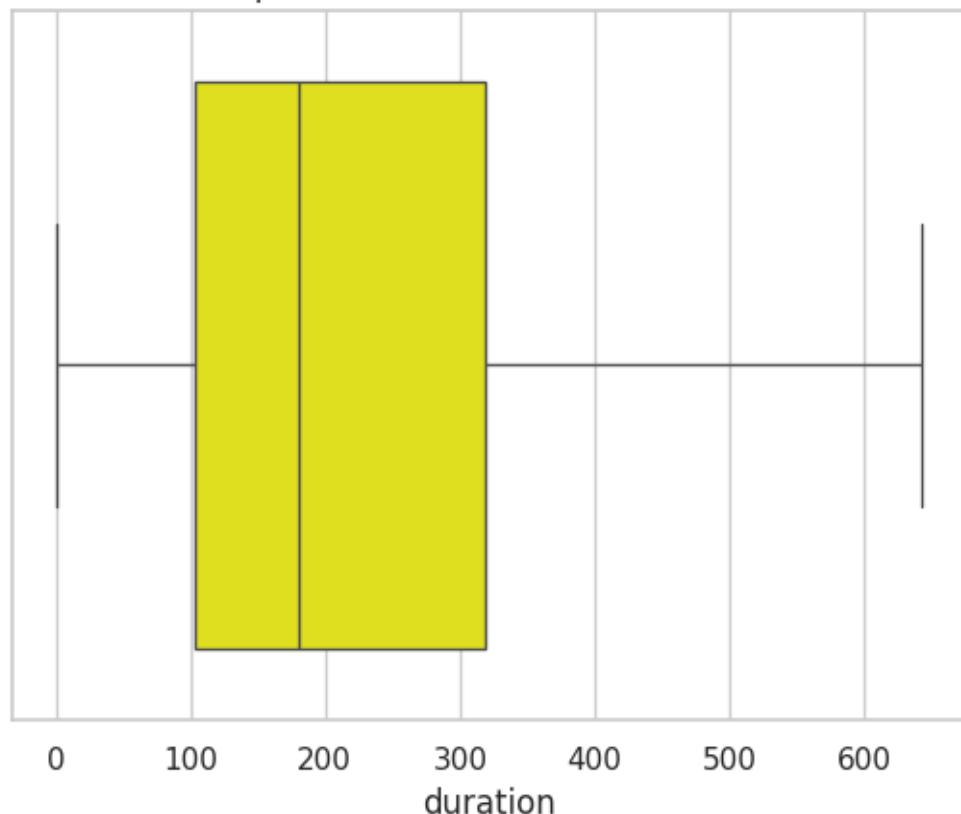


```
[19]: cont_plot(df, 'day' )
```



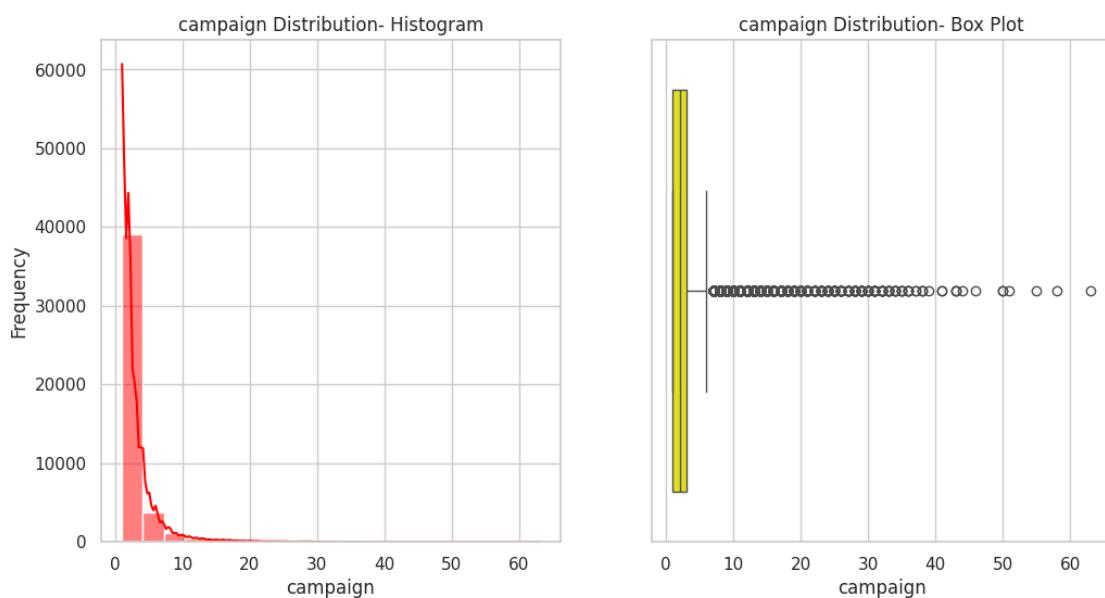
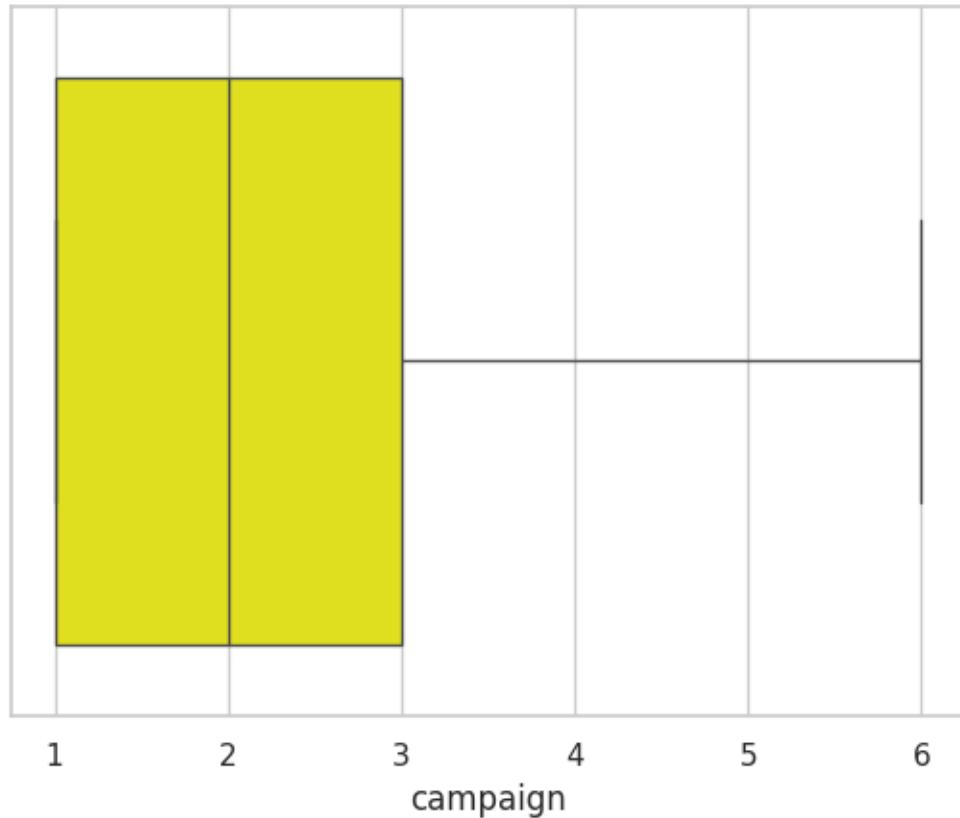
```
[20]: NoOutlier(df, "duration")
cont_plot(df, 'duration' )
```

Boxplot of duration (without outliers)




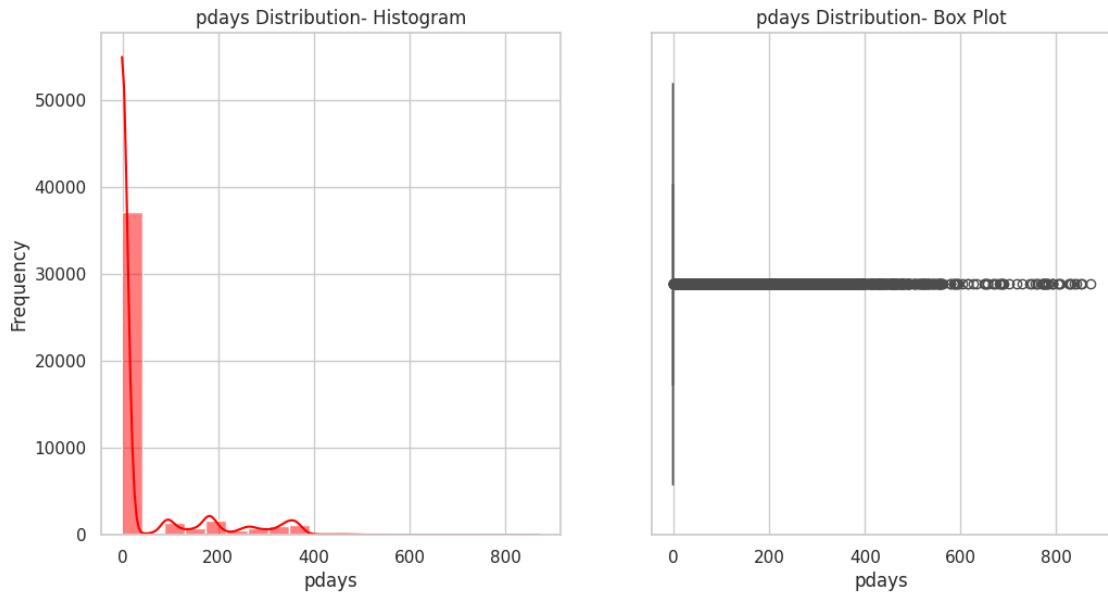
```
[21]: NoOutlier(df , "campaign")
cont_plot(df , 'campaign' )
```

Boxplot of campaign (without outliers)




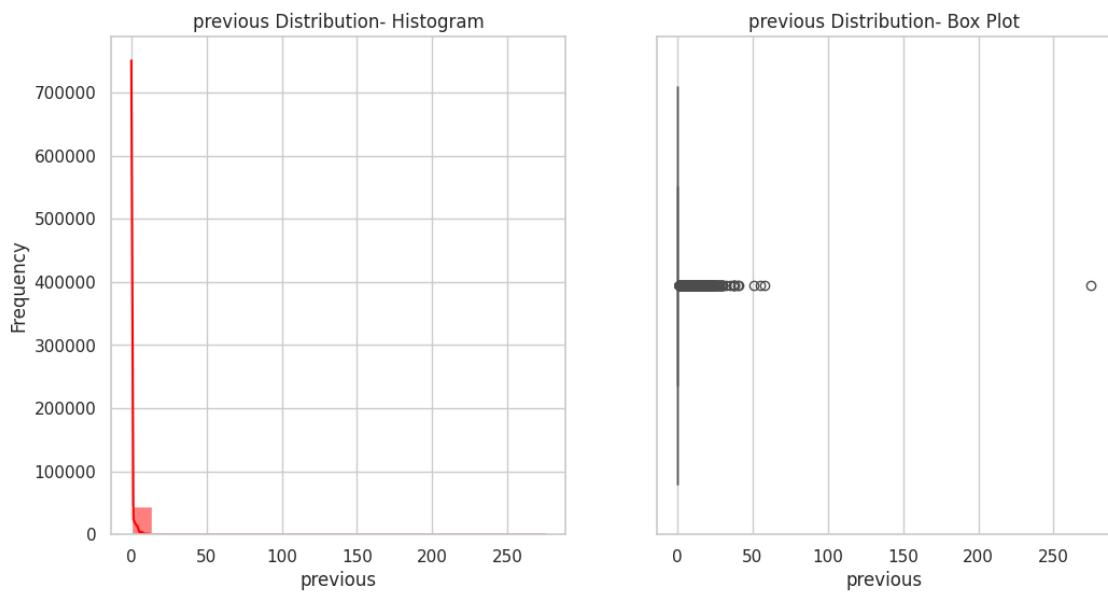
```
[22]: NoOutlier(df,"pdays")
cont_plot(df,'pdays' )
```






```
[23]: NoOutlier(df, "previous")
cont_plot(df, 'previous')
```






```
[24]: # Pairwise scatter plots for numerical variables
sns.pairplot(df.select_dtypes(include=['number']))
plt.show()
```



```
[25]: # Compute correlation matrix
corr_matrix = df.select_dtypes(include=['number']).corr()

# Plot heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0)
plt.title("Pairwise Correlation of Numerical Variables")
plt.show()
```



```
[26]: int_columns = df.select_dtypes('int64').columns.tolist()
print(int_columns)
```

```
['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous']
```

```
[ ]:
```

```
[27]: object_columns = df.select_dtypes(include=['object']).columns.tolist()
```

```
print(object_columns)
```

```
['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
'month', 'poutcome', 'y']
```


```
[28]: def cat_cont_plot(df, xvar, yvar):
    """
    Creates histograms of a continuous variable across categories.

    Parameters:
    - df: pandas DataFrame
    - xvar: str, categorical column name (e.g. 'Target')
    - yvar: str, continuous column name (e.g. 'Age')
    """
    sns.boxplot(x=xvar, y=yvar, data=df, color="pink")
    plt.xticks(rotation=90)


    # Create a larger facet grid of histograms by category
    #g = sns.FacetGrid(df, col=xvar, col_wrap=4, height=6, aspect=0.5) # ↴ Increased size and layout
    #g.map(sns.histplot, yvar, bins=20, color="green")

    # Add titles and labels
    #g.set_axis_labels(yvar, "Frequency")
    #g.set_titles(col_template="{col_name}")
    #plt.tight_layout()
    plt.show()
```

```
[29]: cat_cont_plot(df, "y", "duration")
```



```
[30]: cat_cont_plot(df, "y", "campaign")
```



```
[31]: def cat_cat_plot(dataframe, column_name, hue_column):
    """
    Plots a grouped bar chart showing category frequencies split by hue,
    with both frequency (inside bar) and proportion (above bar) labels.

    Parameters:
    - dataframe: pandas DataFrame
    - column_name: str, name of the categorical column to visualize (x-axis)
    - hue_column: str, name of the second categorical variable to group by (hue)
    """
    # Count combinations of column and hue
    counts_df = dataframe.groupby([column_name, hue_column]).size().
    ↪reset_index(name='count')
    total_counts = dataframe[column_name].value_counts()

    # Set plot style
    sns.set(style="whitegrid")
    plt.figure(figsize=(14, 8)) # Larger frame

    # Create bar plot
```


```

ax = sns.barplot(x=column_name, y='count', hue=hue_column, data=counts_df, palette="muted")
# Annotate bars
for container in ax.containers:
    for bar in container:
        height = bar.get_height()
        x = bar.get_x() + bar.get_width() / 2
        category = bar.get_label()
        base_x = int(round(x)) # used for proportion lookup
        if height > 0:
            ax.text(x, height * 0.5, f'{int(height)}', ha='center', va='center',
                     fontsize=10, color='black', fontweight='bold')
            ax.text(x, height + max(counts_df['count']) * 0.02, f'{height / sum(container.datavalues):.1%}',
                     ha='center', fontsize=10, color='navy', fontweight='bold')


# Beautify plot
plt.title(f'{column_name} Distribution by {hue_column}', fontsize=16)
plt.xlabel(column_name, fontsize=23)
plt.ylabel('Frequency', fontsize=19)
plt.xticks(rotation=45)
plt.legend(title=hue_column)
plt.tight_layout()
plt.show()

```

[32]: cat_cat_plot(df, 'job', "y")



```
[33]: cat_cat_plot(df,'education','y')
```



```
[ ]:
```

```
[34]: #import numpy as np

# Indicator: whether client was contacted before
df['previous_contact'] = (df['pdays'] != -1).astype(int)

# Replace -1 with NaN so pdays is only meaningful when contact exists
df['pdays'] = df['pdays'].replace(-1, np.nan)
```

1 Research Question 1: Which customer and campaign features best predict term deposit subscription?

```
[35]: rq1_vars=["age", "job", "education", "marital", "balance", "default",
             "housing", "loan", "contact", "duration", "poutcome", "previous",
             "campaign", "previous_contact", 'pdays', "y"]
```

```
[36]: dfrq1=df[rq1_vars]
```

```
[37]: dfrq1.head()
```

	age	job	education	marital	balance	default	housing	loan	\
0	58	management	tertiary	married	2143	no	yes	no	
1	44	technician	secondary	single	29	no	yes	no	
2	33	entrepreneur	secondary	married	2	no	yes	yes	
3	47	blue-collar	unknown	married	1506	no	yes	no	
4	33	unknown	unknown	single	1	no	no	no	

	contact	duration	poutcome	previous	campaign	previous_contact	pdays	y
0	unknown	261	unknown	0	1		0	NaN no
1	unknown	151	unknown	0	1		0	NaN no
2	unknown	76	unknown	0	1		0	NaN no
3	unknown	92	unknown	0	1		0	NaN no
4	unknown	198	unknown	0	1		0	NaN no

[]:

[]:

[]:

1.1 Preprocessing

```
[38]: dfrq1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
```

```

Data columns (total 16 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               45211 non-null   int64  
 1   job               45211 non-null   object  
 2   education         45211 non-null   object  
 3   marital           45211 non-null   object  
 4   balance           45211 non-null   int64  
 5   default           45211 non-null   object  
 6   housing            45211 non-null   object  
 7   loan               45211 non-null   object  
 8   contact            45211 non-null   object  
 9   duration           45211 non-null   int64  
 10  poutcome          45211 non-null   object  
 11  previous           45211 non-null   int64  
 12  campaign           45211 non-null   int64  
 13  previous_contact  45211 non-null   int64  
 14  pdays              8257 non-null    float64 
 15  y                  45211 non-null   object  
dtypes: float64(1), int64(6), object(9)
memory usage: 5.5+ MB

```

[]:

[]:

[39]: dfrq1.shape

[39]: (45211, 16)

1.1.1 Encoding all object variables into dummy variable.

```

[40]: cat_cols = dfrq1.select_dtypes(include='object').columns

dfrq1_enc = pd.get_dummies(dfrq1, columns=cat_cols, drop_first=True)

print(dfrq1_enc.shape)
print(dfrq1_enc.head())

```

```
(45211, 32)
   age  balance  duration  previous  campaign  previous_contact  pdays \
0   58      2143     261        0        1                0       NaN
1   44        29     151        0        1                0       NaN
2   33        2      76        0        1                0       NaN
3   47     1506     92        0        1                0       NaN
4   33        1     198        0        1                0       NaN
```

```

job_blue-collar  job_entrepreneur  job_housemaid ... marital_single \
0                  0                 0                 0   ...
1                  0                 0                 0   ...
2                  0                 1                 0   ...
3                  1                 0                 0   ...
4                  0                 0                 0   ...

default_yes  housing_yes  loan_yes  contact_telephone  contact_unknown \
0            0           1         0                   0           1
1            0           1         0                   0           1
2            0           1         1                   0           1
3            0           1         0                   0           1
4            0           0         0                   0           1

poutcome_other  poutcome_success  poutcome_unknown  y_yes
0              0                 0             1       0
1              0                 0             1       0
2              0                 0             1       0
3              0                 0             1       0
4              0                 0             1       0

```

[5 rows x 32 columns]

```
[41]: X = dfrq1_enc.drop(columns=['y_yes'])
y = dfrq1_enc['y_yes']
```

```
[42]: #Train-test split (stratified)
X_train, X_test, y_train, y_test = train_test_split( X, y,
    test_size=0.3,
    stratify=y,
    random_state=42
)
```

Given the presence of substantial outliers in several numerical variables, RobustScaler was used to scale the features, as it relies on the median and interquartile range and is less sensitive to extreme values.

[]:

[]:

[]:

Since pdays is undefined for clients never previously contacted, a binary indicator was introduced to represent prior contact status. Remaining missing values in pdays were imputed using the median to avoid introducing artificial extremes.

The variable pdays uses the value -1 to indicate clients who were never previously contacted. As this value is not a valid numerical quantity, it was replaced with missing values, and a binary

indicator variable (previous_contact) was created to capture prior contact status. Since logistic regression does not handle missing values natively, a median imputation strategy was applied within a modeling pipeline prior to robust scaling and model fitting. RobustScaler was used due to the presence of substantial outliers in several numerical variables.

```
[43]: log_reg_pipeline = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='median')),      # handles NaNs (pdays)  
    ('scaler', RobustScaler()),                          # robust to outliers  
    ('model', LogisticRegression(  
        max_iter=1000,  
        class_weight='balanced',  
        random_state=42  
    ))  
])  
log_reg_pipeline.fit(X_train, y_train)  
y_pred = log_reg_pipeline.predict(X_test)  
  
cm = confusion_matrix(y_test, y_pred)  
print("Confusion matrix of the Logistic Regression \n",cm)  
print("Classification report of the Logistic Regression\n",  
     classification_report(y_test, y_pred))  
#ROC-AUC  
y_prob = log_reg_pipeline.predict_proba(X_test)[:, 1]  
roc_auc = roc_auc_score(y_test, y_prob)  
  
print("ROC-AUC of the Logistic Regression:", roc_auc)
```

```
Confusion matrix of the Logistic Regression  
[[10005 1972]  
 [ 343 1244]]  
Classification report of the Logistic Regression  
             precision    recall  f1-score   support  
  
          0       0.97     0.84     0.90    11977  
          1       0.39     0.78     0.52     1587  
  
    accuracy                           0.83    13564  
   macro avg       0.68     0.81     0.71    13564  
weighted avg       0.90     0.83     0.85    13564
```

ROC-AUC of the Logistic Regression: 0.8917974426830168

Interpret coefficients

```
[44]: import pandas as pd  
import numpy as np  
  
coef = log_reg_pipeline.named_steps['model'].coef_[0]
```

```

coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': coef,
    'Odds_Ratio': np.exp(coef)
}).sort_values(by='Odds_Ratio', ascending=False)

coef_df

```

[44]:

	Feature	Coefficient	Odds_Ratio
29	poutcome_success	2.387607	10.887405
2	duration	1.192009	3.293691
14	job_student	0.564808	1.759110
19	education_tertiary	0.513514	1.671154
11	job_retired	0.429690	1.536781
20	education_unknown	0.305059	1.356705
18	education_secondary	0.258492	1.294975
28	poutcome_other	0.251713	1.286227
22	marital_single	0.187751	1.206533
5	previous_contact	0.182628	1.200368
26	contact_telephone	0.053333	1.054781
1	balance	0.040106	1.040921
3	previous	0.030486	1.030956
0	age	0.007305	1.007332
6	pdays	0.000238	1.000239
30	poutcome_unknown	-0.098979	0.905762
21	marital_married	-0.180855	0.834557
10	job_management	-0.204880	0.814745
4	campaign	-0.235807	0.789933
23	default_yes	-0.251401	0.777711
16	job_unemployed	-0.282399	0.753973
15	job_technician	-0.326132	0.721710
17	job_unknown	-0.366712	0.693009
13	job_services	-0.444485	0.641155
8	job_entrepreneur	-0.536855	0.584584
9	job_housemaid	-0.543343	0.580803
7	job_blue-collar	-0.543642	0.580630
12	job_self-employed	-0.545387	0.579617
25	loan_yes	-0.685377	0.503900
24	housing_yes	-0.849504	0.427627
27	contact_unknown	-1.326366	0.265440

[]:

Interpretation rule:

Odds Ratio > 1 → increases probability of subscription

Odds Ratio < 1 → decreases probability


```
[45]: #Extract coefficients and odds ratios from the pipeline
coef = log_reg_pipeline.named_steps['model'].coef_[0]

coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': coef,
    'Odds_Ratio': np.exp(coef)
})

coef_df = coef_df.sort_values(by='Odds_Ratio', ascending=False)
```

```
[46]: #Visualise coefficients (direction + strength)
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 10))
plt.barh(coef_df['Feature'], coef_df['Coefficient'])
plt.axvline(0)
plt.xlabel('Logistic Regression Coefficient')
plt.title('Logistic Regression Coefficients')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```




```
[47]: #Visualise odds ratios
plt.figure(figsize=(8, 10))
plt.barh(coef_df['Feature'], coef_df['Odds_Ratio'])
plt.axvline(1)
plt.xlabel('Odds Ratio')
```

```

plt.title('Odds Ratios from Logistic Regression')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```



```
# Fit Random Forest (no scaling needed)
```

Random Forest does NOT need scaling and handles outliers naturally.

```
[48]: rf = RandomForestClassifier(  
        n_estimators=300,  
        max_depth=None,  
        min_samples_leaf=5,  
        class_weight='balanced',  
        random_state=42,  
        n_jobs=-1  
)  
  
rf.fit(X_train, y_train)
```

```
[48]: RandomForestClassifier(class_weight='balanced', min_samples_leaf=5,  
                           n_estimators=300, n_jobs=-1, random_state=42)
```

Model evaluation

```
[49]: y_pred = rf.predict(X_test)  
y_prob = rf.predict_proba(X_test)[:, 1]  
  
print("The Confusion Matrix of Random Forest:\n", confusion_matrix(y_test, y_pred))  
print("\nClassification Report of Random Forest:\n", classification_report(y_test, y_pred))  
print("The ROC-AUC of Random Forest:", roc_auc_score(y_test, y_prob))
```

The Confusion Matrix of Random Forest:

```
[[10483 1494]  
 [ 386 1201]]
```

Classification Report of Random Forest:

	precision	recall	f1-score	support
0	0.96	0.88	0.92	11977
1	0.45	0.76	0.56	1587
accuracy			0.86	13564
macro avg	0.71	0.82	0.74	13564
weighted avg	0.90	0.86	0.88	13564

The ROC-AUC of Random Forest: 0.9023631672951818

Feature importance

```
[50]: feature_importance = pd.DataFrame({  
    'Feature': X.columns,
```

```
'Importance': rf.feature_importances_
}).sort_values(by='Importance', ascending=False)


feature_importance.head(15)
```

```
[50]:
```

	Feature	Importance
2	duration	0.469407
1	balance	0.080987
0	age	0.077718
27	contact_unknown	0.055870
6	pdays	0.047498
29	poutcome_success	0.047468
24	housing_yes	0.041986
4	campaign	0.035903
3	previous	0.019197
5	previous_contact	0.012006
30	poutcome_unknown	0.011736
25	loan_yes	0.011049
21	marital_married	0.010674
7	job_blue-collar	0.009519
19	education_tertiary	0.008935

```
[51]: #Visualise top feature importances
top_features = feature_importance.head(15)

plt.figure(figsize=(8, 8))
plt.barh(top_features['Feature'], top_features['Importance'])
plt.xlabel('Feature Importance')
plt.title('Top 15 Feature Importances (Random Forest)')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```



A Random Forest classifier was trained to identify key customer and campaign features associated with term deposit subscription. As a tree-based model, Random Forest does not require feature scaling and is robust to outliers. Class imbalance was addressed using balanced class weights. Model performance was evaluated using recall, F1-score, and ROC–AUC. Feature importance scores were extracted to identify the most influential predictors.

1.2 SHAP

To enhance interpretability of the Random Forest model, SHAP (SHapley Additive exPlanations) values were computed. SHAP provides a model-agnostic explanation framework that quantifies both the magnitude and direction of each feature's contribution to the predicted outcome. The SHAP summary plot highlights the most influential features and reveals whether higher or lower values of each feature increase the likelihood of subscription.

```
[52]: # Take a random sample of the test set
X_test_shap = X_test.sample(n=1000, random_state=42)
```


```
[53]: #explainer = shap.TreeExplainer(rf)
#shap_values = explainer.shap_values(X_test_shap)
```

```
[54]: # Use the unified SHAP API
explainer = shap.Explainer(rf, X_train)

# Compute SHAP values on the sample
shap_values = explainer(X_test_shap)

# Summary plot (this WILL work)
shap.summary_plot(
    shap_values.values,
    X_test_shap,
    plot_type="dot",
    show=True
)
```

100%|=====| 1996/2000 [02:34<00:00]



```
[ ]:
```

```
[ ]:
```

```
[55]: # -----
```

```
# Compute SHAP INTERACTION values (slow but needed once)
```

```
# -----
```

```
shap_interaction_values = explainer.shap_interaction_values(X_test_shap)
```

```

# For binary classification: take positive class ("yes")
shap_inter_pos = shap_interaction_values[:, :, 1]

# -----
# Quantify MAIN vs INTERACTION effects
# -----
# ---- MAIN EFFECTS (positive class only) ----
shap_main = np.mean(
    np.abs(shap_values.values[:, :, 1]), # class "yes"
    axis=0
)
# Mean absolute MAIN effects per feature
#shap_main = np.mean(np.abs(shap_values.values), axis=0)

# Mean absolute INTERACTION effects
shap_inter = np.mean(np.abs(shap_inter_pos), axis=0)

# Remove self-interactions (diagonal)
np.fill_diagonal(shap_inter, 0)

# Aggregate interaction strength per feature
interaction_strength = shap_inter.sum(axis=1)

# -----
# Comparison table (proof of weak interactions)
# -----
comparison = pd.DataFrame({
    "feature": X_test_shap.columns,
    "main_effect": shap_main,
    "interaction_effect": interaction_strength,
    "interaction_ratio": interaction_strength / shap_main
}).sort_values("interaction_ratio", ascending=False)

comparison

```

FEATURE_DEPENDENCE::independent does not support interactions!

[55]:	feature	main_effect	interaction_effect	interaction_ratio
0	age	0.016881	0.0	0.0
16	job_unemployed	0.000485	0.0	0.0
29	poutcome_success	0.016465	0.0	0.0
28	poutcome_other	0.000807	0.0	0.0
27	contact_unknown	0.046133	0.0	0.0
26	contact_telephone	0.000857	0.0	0.0
25	loan_yes	0.009922	0.0	0.0
24	housing_yes	0.037571	0.0	0.0

23	default_yes	0.000445	0.0	0.0
22	marital_single	0.007459	0.0	0.0
21	marital_married	0.010193	0.0	0.0
20	education_unknown	0.000677	0.0	0.0
19	education_ternary	0.013163	0.0	0.0
18	education_secondary	0.002640	0.0	0.0
17	job_unknown	0.000004	0.0	0.0
15	job_technician	0.002213	0.0	0.0
1	balance	0.020867	0.0	0.0
14	job_student	0.001402	0.0	0.0
13	job_services	0.001417	0.0	0.0
12	job_self-employed	0.000555	0.0	0.0
11	job_retired	0.000969	0.0	0.0
10	job_management	0.003568	0.0	0.0
9	job_housemaid	0.000505	0.0	0.0
8	job_entrepreneur	0.000599	0.0	0.0
7	job_blue-collar	0.005405	0.0	0.0
6	pdays	0.012604	0.0	0.0
5	previous_contact	0.007716	0.0	0.0
4	campaign	0.017559	0.0	0.0
3	previous	0.009495	0.0	0.0
2	duration	0.134146	0.0	0.0
30	poutcome_unknown	0.007265	0.0	0.0

[]:

Although Random Forest models are capable of capturing nonlinear interactions, SHAP interaction analysis shows that interaction effects are negligible for all predictors. The model's predictions are therefore driven primarily by additive contributions of individual features rather than by complex feature interactions.

Methods / Explainability section

SHAP interaction values were computed on a representative subset of the test data to assess whether the Random Forest model relied on nonlinear feature interactions. Interaction effects were quantified and compared to main feature effects using the ratio of interaction strength to main effect magnitude.

Results / Interpretation section

Across all predictors, interaction effects were effectively zero relative to main effects. This indicates that the Random Forest model predominantly relies on additive feature contributions, with no evidence of strong pairwise interactions influencing subscription predictions.

This finding suggests that the predictive structure of the data is largely linear-additive, despite the use of a non-linear model.

2 RQ2 Can machine learning models outperform traditional statistical models in predicting subscription?

2.0.1 Logistic Regression

```
[56]: log_reg_pipeline = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='median')),      # handles NaNs (pdays)  
    ('scaler', RobustScaler()),                          # robust to outliers  
    ('model', LogisticRegression(  
        max_iter=1000,  
        class_weight='balanced',  
        random_state=42  
    ))  
])  
log_reg_pipeline.fit(X_train, y_train)
```

```
[56]: Pipeline(steps=[('imputer', SimpleImputer(strategy='median')),  
                     ('scaler', RobustScaler()),  
                     ('model',  
                      LogisticRegression(class_weight='balanced', max_iter=1000,  
                                         random_state=42))])
```

```
[57]: y_pred = log_reg_pipeline.predict(X_test)  
  
cm = confusion_matrix(y_test, y_pred)  
#ROC-AUC  
y_prob = log_reg_pipeline.predict_proba(X_test)[:, 1]  
roc_auc = roc_auc_score(y_test, y_prob)  
  
print("The ROC-AUC of Logistic Regression:", roc_auc)  
print("\n the confusion matrix of Logistic regression:\n", cm)  
print("\n the classification report of Logistic Regression\n",  
classification_report(y_test, y_pred))
```

The ROC-AUC of Logistic Regression: 0.8917974426830168

```
the confusion matrix of Logistic regression:  
[[10005  1972]  
 [ 343  1244]]  
  
the classification report of Logistic Regression  
precision    recall   f1-score   support  
  
          0       0.97     0.84     0.90     11977  
          1       0.39     0.78     0.52     1587  
  
accuracy                           0.83     13564  
macro avg       0.68     0.81     0.71     13564
```

```
weighted avg      0.90      0.83      0.85      13564
```

```
[58]: import pandas as pd
import numpy as np

coef = log_reg_pipeline.named_steps['model'].coef_[0]

coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': coef,
    'Odds_Ratio': np.exp(coef)
}).sort_values(by='Odds_Ratio', ascending=False)

coef_df
```

```
[58]:          Feature   Coefficient   Odds_Ratio
29     poutcome_success      2.387607      10.887405
2           duration      1.192009      3.293691
14      job_student      0.564808      1.759110
19  education_tertiary      0.513514      1.671154
11      job_retired      0.429690      1.536781
20  education_unknown      0.305059      1.356705
18  education_secondary      0.258492      1.294975
28     poutcome_other      0.251713      1.286227
22      marital_single      0.187751      1.206533
5    previous_contact      0.182628      1.200368
26  contact_telephone      0.053333      1.054781
1        balance      0.040106      1.040921
3        previous      0.030486      1.030956
0         age      0.007305      1.007332
6        pdays      0.000238      1.000239
30    poutcome_unknown     -0.098979      0.905762
21      marital_married     -0.180855      0.834557
10      job_management     -0.204880      0.814745
4        campaign     -0.235807      0.789933
23      default_yes     -0.251401      0.777711
16      job_unemployed     -0.282399      0.753973
15      job_technician     -0.326132      0.721710
17      job_unknown     -0.366712      0.693009
13      job_services     -0.444485      0.641155
8    job_entrepreneur     -0.536855      0.584584
9      job_housemaid     -0.543343      0.580803
7    job_blue-collar     -0.543642      0.580630
12  job_self-employed     -0.545387      0.579617
25      loan_yes     -0.685377      0.503900
24      housing_yes     -0.849504      0.427627
```


```
27      contact_unknown    -1.326366    0.265440
```

```
[59]: #Extract coefficients and odds ratios from the pipeline  
coef = log_reg_pipeline.named_steps['model'].coef_[0]
```

```
coef_df = pd.DataFrame({  
    'Feature': X.columns,  
    'Coefficient': coef,  
    'Odds_Ratio': np.exp(coef)  
})  
  
coef_df = coef_df.sort_values(by='Odds_Ratio', ascending=False)
```

```
[60]: #Visualise coefficients (direction + strength)  
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(8, 10))  
plt.barh(coef_df['Feature'], coef_df['Coefficient'])  
plt.axvline(0)  
plt.xlabel('Logistic Regression Coefficient')  
plt.title('Logistic Regression Coefficients')  
plt.gca().invert_yaxis()  
plt.tight_layout()  
plt.show()
```




```
[61]: #Visualise odds ratios
plt.figure(figsize=(8, 10))
plt.barh(coef_df['Feature'], coef_df['Odds_Ratio'])
plt.axvline(1)
plt.xlabel('Odds Ratio')
```

```

plt.title('Odds Ratios from Logistic Regression')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```



2.0.2 Decision Tree

```
[62]: #model_dt = DecisionTreeClassifier(random_state=42)
model_dt = DecisionTreeClassifier(
    max_depth=4,          # limit depth
    min_samples_split=5,
    min_samples_leaf=2,
    random_state=42
)
model_dt.fit(X_train, y_train)
y_pred_dt = model_dt.predict(X_test)
# Get predicted probabilities for the positive class
y_prob_dt = model_dt.predict_proba(X_test)[:, 1]

# Compute AUC
auc_dt = roc_auc_score(y_test, y_prob_dt)

print("The AUC of Decision Tree:", auc_dt)
print("Accuracy of Decision Tree:", accuracy_score(y_test, y_pred_dt))
print("\nClassification Report of Decision Tree:\n",classification_report(y_test, y_pred_dt))
cm_dt = confusion_matrix(y_test, y_pred_dt)
print("\nConfusion Matrix of Decision Tree:\n", cm_dt)
```

The AUC of Decision Tree: 0.8336726204746873

Accuracy of Decision Tree: 0.898849896785609

Classification Report of Decision Tree:


	precision	recall	f1-score	support
0	0.91	0.98	0.94	11977
1	0.65	0.29	0.40	1587
accuracy			0.90	13564
macro avg	0.78	0.64	0.67	13564
weighted avg	0.88	0.90	0.88	13564

Confusion Matrix of Decision Tree:

```
[[11730  247]
 [ 1125  462]]
```


```
[63]: plt.figure(figsize=(6,5))
sns.heatmap(cm_dt, annot=True, fmt='d', cmap='Blues',
            xticklabels=model_dt.classes_,
            yticklabels=model_dt.classes_)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
```

```
plt.title("Confusion Matrix - Decision Tree")
plt.show()
```



```
[64]: fpr, tpr, thresholds = roc_curve(y_test, y_prob_dt)


plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f"Decision Tree (AUC = {auc_dt:.3f})")
plt.plot([0, 1], [0, 1], linestyle="--") # random model line
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Decision Tree")
plt.legend()
plt.show()
```



```
[65]: plt.figure(figsize=(12,8))

tree.plot_tree(
    model_dt,
    feature_names=X_train.columns,           # <-- use dataframe columns
    class_names=model_dt.classes_.astype(str), # <-- get class labels
    filled=True
)

plt.show()
```



2.0.3 Random Forest

```
[66]: rf2 = RandomForestClassifier(  
        n_estimators=300,  
        max_depth=None,  
        min_samples_leaf=5,  
        class_weight='balanced',  
        random_state=42,  
        n_jobs=-1  
)  
  
rf2.fit(X_train, y_train)  
y_pred_rf = rf2.predict(X_test)  
y_prob_rf = rf2.predict_proba(X_test)[:, 1]  
  
print("The confusion matrix of Random Forest:\n", confusion_matrix(y_test, y_pred_rf))  
print("\nThe classification Report of Random Forest:\n", classification_report(y_test, y_pred_rf))  
print("The ROC-AUC of Random Forest:", roc_auc_score(y_test, y_prob_rf))
```

The confusion matrix of Random Forest:

`[[10483 1494]`

```
[ 386 1201]]
```

The classification Report of Random Forest:

	precision	recall	f1-score	support
0	0.96	0.88	0.92	11977
1	0.45	0.76	0.56	1587
accuracy			0.86	13564
macro avg	0.71	0.82	0.74	13564
weighted avg	0.90	0.86	0.88	13564

The ROC-AUC of Random Forest: 0.9023631672951818


```
[67]: # Compute ROC curve
fpr_rf, tpr_rf, thresholds = roc_curve(y_test, y_prob_rf)

# Compute AUC (again, for label in plot)
auc_rf = roc_auc_score(y_test, y_prob_rf)

# Plot
plt.figure(figsize=(8,6))
plt.plot(fpr_rf, tpr_rf, linewidth=2,
         label=f"Random Forest (AUC = {auc_rf:.3f})")

# Diagonal baseline
plt.plot([0,1], [0,1], linestyle='--', linewidth=1)

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Random Forest")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[68]: feature_importance_rf = pd.DataFrame({
    'Feature': X.columns,
    'Importance': rf2.feature_importances_
}).sort_values(by='Importance', ascending=False)

feature_importance_rf.head(15)
```

	Feature	Importance
2	duration	0.469407
1	balance	0.080987
0	age	0.077718
27	contact_unknown	0.055870
6	pdays	0.047498
29	poutcome_success	0.047468
24	housing_yes	0.041986
4	campaign	0.035903
3	previous	0.019197
5	previous_contact	0.012006
30	poutcome_unknown	0.011736
25	loan_yes	0.011049
21	marital_married	0.010674


```

7      job_blue-collar    0.009519
19   education_tertiary   0.008935

```

```
[69]: #Visualise top feature importances
top_features_rf = feature_importance_rf.head(15)

plt.figure(figsize=(8, 8))
plt.barh(top_features_rf['Feature'], top_features_rf['Importance'])
plt.xlabel('Feature Importance')
plt.title('Top 15 Feature Importances (Random Forest)')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```



2.0.4 Extreme Gradient Boosting

```
[70]: model_xgb = XGBClassifier(  
    n_estimators=200,  
    learning_rate=0.05,  
    max_depth=4,  
    subsample=0.8,  
    colsample_bytree=0.8,  
    tree_method="hist",  
    random_state=42,  
    eval_metric="logloss"  
)  
  
# Fit with NumPy  
model_xgb.fit(  
    X_train.to_numpy(dtype=np.float32),  
    y_train.to_numpy(dtype=np.float32)  
)  
  
# Predict with NumPy  
y_pred_xgb = model_xgb.predict(  
    X_test.to_numpy(dtype=np.float32)  
)  
  
print("Accuracy of xgboost:", accuracy_score(y_test, y_pred_xgb))  
print("\nClassification Report of xgboost:\n", classification_report(y_test, y_pred_xgb))  
print("\nConfusion Matrix of xgboost:\n", confusion_matrix(y_test, y_pred_xgb))
```

Accuracy of xgboost: 0.9030521969920378

Classification Report of xgboost:

	precision	recall	f1-score	support
0	0.92	0.97	0.95	11977
1	0.64	0.39	0.48	1587
accuracy			0.90	13564
macro avg	0.78	0.68	0.71	13564
weighted avg	0.89	0.90	0.89	13564

Confusion Matrix of xgboost:

```
[[11634  343]  
 [ 972  615]]
```

```
[71]: # Get predicted probabilities (positive class)
y_prob_xgb = model_xgb.predict_proba(
    X_test.to_numpy(dtype=np.float32)
)[:, 1]

# Compute AUC
auc_xgb = roc_auc_score(y_test, y_prob_xgb)
print("ROC-AUC (XGBoost):", auc_xgb)
```


ROC-AUC (XGBoost): 0.907066784535935

```
[72]: # Compute ROC curve
fpr_xgb, tpr_xgb, thresholds = roc_curve(y_test, y_prob_xgb)

plt.figure(figsize=(8,6))
plt.plot(fpr_xgb, tpr_xgb, linewidth=2,
         label=f"XGBoost (AUC = {auc_xgb:.3f})")

# Random baseline
plt.plot([0,1], [0,1], linestyle='--')

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - XGBoost")
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```



[]:

```
[73]: # Save feature names BEFORE conversion
feature_names = X_train.columns.tolist()

# Convert to numpy
X_train_np = X_train.to_numpy(dtype=np.float32)
X_test_np = X_test.to_numpy(dtype=np.float32)
y_train_np = y_train.to_numpy(dtype=np.float32)


# Train
model_xgb.fit(X_train_np, y_train_np)

# After training, manually set feature names
model_xgb.get_booster().feature_names = feature_names
```

```
[74]: # Get importance values
importances = model_xgb.feature_importances_

# Create DataFrame
importance_df = pd.DataFrame({
```

```
"Feature": feature_names,  
    "Importance": importances  
})  
  
# Sort by importance  
importance_df = importance_df.sort_values(by="Importance", ascending=True)  
  
# Plot  
plt.figure(figsize=(12, 16))    # Bigger figure  
plt.barh(importance_df["Feature"], importance_df["Importance"])  
plt.xlabel("Feature Importance")  
plt.title("XGBoost Feature Importance")  
plt.tight_layout()  
plt.show()
```



3 Research Question 3: How early in a campaign can reliable predictions be made?

3.0.1 Scenario 1: the early-stage, pre-campaign or start of the campaign

[]:

```
[75]: rq3_vars_sc1=["age", "job","education", "marital", "balance", "default",  
    ↪"housing","loan","y"]  
dfrq3_sc1=df[rq3_vars_sc1]  
#Encoding all object variables into dummy variable.  
cat_cols_sc1 = dfrq3_sc1.select_dtypes(include='object').columns  
  
dfrq3_sc1_enc = pd.get_dummies(dfrq3_sc1, columns=cat_cols_sc1, drop_first=True)  
  
print(dfrq3_sc1_enc.shape)  
#print(dfrq3_sc1_enc.head())  
  
#Independent and depend variables split  
X_sc1 = dfrq3_sc1_enc.drop(columns=['y_yes'])  
y_sc1 = dfrq3_sc1_enc['y_yes']  
#Train-test split (stratified)  
#from sklearn.model_selection import train_test_split  
  
X_train_sc1, X_test_sc1, y_train_sc1, y_test_sc1 = train_test_split( X_sc1,  
    ↪y_sc1,  
    test_size=0.3,  
    stratify=y_sc1,  
    random_state=42  
)  
  
log_reg_pipeline_sc1 = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='median')), # handles NaNs (pdays)  
    ('scaler', RobustScaler()), # robust to outliers  
    ('model', LogisticRegression(  
        max_iter=1000,  
        class_weight='balanced',  
        random_state=42  
    ))  
])  
log_reg_pipeline_sc1.fit(X_train_sc1, y_train_sc1)  
  
#model evaluation  
y_pred_sc1 = log_reg_pipeline_sc1.predict(X_test_sc1)  
  
cm_sc1 = confusion_matrix(y_test_sc1, y_pred_sc1)  
print("The confusion matrix of Scenario 1 is : \n",cm_sc1)
```

```

print("The classification report of Scenario 1 is :")
    ↵\n",classification_report(y_test_sc1, y_pred_sc1))

#ROC-AUC
#from sklearn.metrics import roc_auc_score

y_prob_sc1 = log_reg_pipeline_sc1.predict_proba(X_test_sc1)[:, 1]
roc_auc_sc1 = roc_auc_score(y_test_sc1, y_prob_sc1)

print("The ROC-AUC of Scenario 1 is:", roc_auc_sc1)

```

(45211, 22)

The confusion matrix of Scenario 1 is :

```

[[7543 4434]
 [ 597  990]]

```

The classification report of Scenario 1 is :

	precision	recall	f1-score	support
0	0.93	0.63	0.75	11977
1	0.18	0.62	0.28	1587
accuracy			0.63	13564
macro avg	0.55	0.63	0.52	13564
weighted avg	0.84	0.63	0.70	13564

The ROC-AUC of Scenario 1 is: 0.6689852515578194

3.0.2 Scenario 2 : mid-stage or during the campaign stage

[]:

```

[77]: rq3_sc2_vars=["age", "job", "education", "marital", "balance", "default",
    ↵"housing", "loan", 'previous_contact', 'pdays', "poutcome", "y"]
dfrq3_sc2=df[rq3_sc2_vars]
#Encoding all object variables into dummy variable.
cat_cols_sc2= dfrq3_sc2.select_dtypes(include='object').columns

dfrq3_sc2_enc = pd.get_dummies(dfrq3_sc2, columns=cat_cols_sc2, drop_first=True)

print(dfrq3_sc2_enc.shape)
#print(dfrq3_sc2_enc.head())

#Independent and depend variables split
X_sc2 = dfrq3_sc2_enc.drop(columns=['y_yes'])
y_sc2 = dfrq3_sc2_enc['y_yes']
#Train-test split (stratified)

```

```

X_train_sc2, X_test_sc2, y_train_sc2, y_test_sc2 = train_test_split( X_sc2, ▾
    ↵y_sc2,
    test_size=0.3,
    stratify=y_sc2,
    random_state=42
)

log_reg_pipeline_sc2 = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),    # handles NaNs (pdays)
    ('scaler', RobustScaler()),                         # robust to outliers
    ('model', LogisticRegression(
        max_iter=1000,
        class_weight='balanced',
        random_state=42
    ))
])
log_reg_pipeline_sc2.fit(X_train_sc2, y_train_sc2)

#model evaluation
y_pred_sc2 = log_reg_pipeline_sc2.predict(X_test_sc2)

cm_sc2 = confusion_matrix(y_test_sc2, y_pred_sc2)
print("The confusion matrix of Scenario 2 is: \n",cm_sc2)
print("The classification report of Scenario 2 is: \n"
    ↵\n",classification_report(y_test_sc2, y_pred_sc2))

#ROC-AUC
#from sklearn.metrics import roc_auc_score

y_prob_sc2 = log_reg_pipeline_sc2.predict_proba(X_test_sc2)[:, 1]
roc_auc_sc2 = roc_auc_score(y_test_sc2, y_prob_sc2)

print("The ROC-AUC of Scenario 2 is:", roc_auc_sc2)

```

(45211, 27)

The confusion matrix of Scenario 2 is:

[[8707 3270]
[643 944]]

The classification report of Scenario 2 is:

	precision	recall	f1-score	support
0	0.93	0.73	0.82	11977
1	0.22	0.59	0.33	1587
accuracy			0.71	13564
macro avg	0.58	0.66	0.57	13564

```
weighted avg      0.85      0.71      0.76      13564
```

The ROC-AUC of Scenario 2 is: 0.7218435997287176

3.0.3 Scenario 3 the late stage or active engagement stage

```
[ ]: #rq1_vars=["age", "job", "education", "marital", "balance", "default",  
    ↪ "housing", "loan", "contact", "duration", "poutcome", "previous", "campaign"  
    ↪, 'previous_contact', 'pdays', 'y']  
  
[78]: rq3_sc3_vars=["age", "job", "education", "marital", "balance", "default",  
    ↪ "housing", "loan", 'previous_contact', 'pdays', "poutcome", "duration", "campaign",  
    ↪ "contact", "y"]  
dfrq3_sc3=df[rq3_sc3_vars]  
#Encoding all object variables into dummy variable.  
cat_cols_sc3= dfrq3_sc3.select_dtypes(include='object').columns  
  
dfrq3_sc3_enc = pd.get_dummies(dfrq3_sc3, columns=cat_cols_sc3, drop_first=True)  
  
print(dfrq3_sc3_enc.shape)  
#print(dfrq3_sc2_enc.head())  
  
#Independent and depend variables split  
X_sc3 = dfrq3_sc3_enc.drop(columns=['y_yes'])  
y_sc3 = dfrq3_sc3_enc['y_yes']  
#Train-test split (stratified)  
  
X_train_sc3, X_test_sc3, y_train_sc3, y_test_sc3 = train_test_split( X_sc3,  
    ↪ y_sc3,  
    test_size=0.3,  
    stratify=y_sc3,  
    random_state=42  
)  
  
log_reg_pipeline_sc3 = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='median')),      # handles NaNs (pdays)  
    ('scaler', RobustScaler()),                          # robust to outliers  
    ('model', LogisticRegression(  
        max_iter=1000,  
        class_weight='balanced',  
        random_state=42  
    ))  
)  
log_reg_pipeline_sc3.fit(X_train_sc3, y_train_sc3)  
  
#model evaluation
```

```

y_pred_sc3 = log_reg_pipeline_sc3.predict(X_test_sc3)

cm_sc3 = confusion_matrix(y_test_sc3, y_pred_sc3)
print("The confusion matrix of Scenario 3 is: \n",cm_sc3)
print("The classification report of Scenario 3 is:\n",classification_report(y_test_sc3, y_pred_sc3))

#ROC-AUC
from sklearn.metrics import roc_auc_score

y_prob_sc3 = log_reg_pipeline_sc3.predict_proba(X_test_sc3)[:, 1]
roc_auc_sc3 = roc_auc_score(y_test_sc3, y_prob_sc3)

print("The ROC-AUC of Scenario 3 is:", roc_auc_sc3)

```

(45211, 31)

The confusion matrix of Scenario 3 is:

[[10005	1972]
[340	1247]]

The classification report of Scenario 3 is:

	precision	recall	f1-score	support
0	0.97	0.84	0.90	11977
1	0.39	0.79	0.52	1587
accuracy			0.83	13564
macro avg	0.68	0.81	0.71	13564
weighted avg	0.90	0.83	0.85	13564

The ROC-AUC of Scenario 3 is: 0.8918576294545643

[]:

4 Research Question 4: Are prediction errors biased across demographic groups?

```

[80]: #Define Variables
rq4_vars = [
    "age", "job", "education", "marital", "balance", "default",
    "housing", "loan", "contact", "duration", "poutcome",
    "previous", "campaign", 'previous_contact', 'pdays'
]

X_rq4 = df[rq4_vars]  #X
y_rq4 = df["y"].map({"yes":1, "no":0})
#Train/Test Split

```

```

#from sklearn.model_selection import train_test_split

X_train_rq4, X_test_rq4, y_train_rq4, y_test_rq4 = train_test_split(
    X_rq4, y_rq4,
    test_size=0.3,
    stratify=y_rq4,
    random_state=42
)
#One-Hot Encode Categorical Variables
X_train_enc_rq4 = pd.get_dummies(X_train_rq4, drop_first=True)
X_test_enc_rq4 = pd.get_dummies(X_test_rq4, drop_first=True)

# Align columns to avoid mismatch
X_train_enc_rq4, X_test_enc_rq4 = X_train_enc_rq4.align(
    X_test_enc_rq4,
    join='left',
    axis=1,
    fill_value=0
)

```

[81]: #Logistic Regression Pipeline

```

#from sklearn.pipeline import Pipeline
#from sklearn.impute import SimpleImputer
#from sklearn.preprocessing import RobustScaler
#from sklearn.linear_model import LogisticRegression

log_reg_pipeline_rq4 = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', RobustScaler()),
    ('model', LogisticRegression(
        max_iter=1000,
        class_weight='balanced',
        random_state=42
    ))
])
log_reg_pipeline_rq4.fit(X_train_enc_rq4, y_train_rq4)

y_pred_rq4 = log_reg_pipeline_rq4.predict(X_test_enc_rq4)

#Create Results DataFrame:
#For bias analysis, we need demographic variables from TEST SET.
results_rq4 = X_test_rq4 .copy()
results_rq4 ["y_true"] = y_test_rq4
results_rq4 ["y_pred"] = y_pred_rq4

#Function to Compute FPR and FNR

```

```

#from sklearn.metrics import confusion_matrix
#import pandas as pd

def error_rates_by_group(data, group_var):

    output = []

    for group in data[group_var].unique():

        subset = data[data[group_var] == group]

        if len(subset) > 0:
            tn, fp, fn, tp = confusion_matrix(
                subset["y_true"],
                subset["y_pred"],
                labels=[0,1]
            ).ravel()

            fpr = fp / (fp + tn) if (fp + tn) > 0 else 0
            fnr = fn / (fn + tp) if (fn + tp) > 0 else 0

            output.append([group, fpr, fnr])

    return pd.DataFrame(output, columns=[group_var, "False Positive Rate", "False Negative Rate"])

```

4.0.1 Bias Analysis

```

[82]: #Marital Status
print("\n Bias analysis results for Marital Status\n",
      error_rates_by_group(results_rq4, "marital"))

#Education
print("\n Bias analysis results for Education\n",
      error_rates_by_group(results_rq4, "education"))

#Job
print("\n Bias analysis results for Job\n",
      error_rates_by_group(results_rq4, "job"))

#Age Groups
results_rq4["age_group"] = pd.cut(
    results_rq4["age"],
    bins=[0,30,40,50,60,100],
    labels=["<30", "30-40", "40-50", "50-60", "60+"]
)

```

```

print("\n Bias analysis results for Age groups \n",  

    ↪error_rates_by_group(results_rq4, "age_group"))

```

Bias analysis results for Marital Status

	marital	False Positive Rate	False Negative Rate
0	divorced	0.149286	0.222857
1	single	0.234251	0.171480
2	married	0.136445	0.243590

Bias analysis results for Education

	education	False Positive Rate	False Negative Rate
0	secondary	0.144928	0.256233
1	tertiary	0.230635	0.166392
2	primary	0.096143	0.225989
3	unknown	0.205628	0.209877

Bias analysis results for Job

	job	False Positive Rate	False Negative Rate
0	blue-collar	0.091355	0.304147
1	unemployed	0.211180	0.222222
2	self-employed	0.131336	0.173913
3	services	0.106074	0.320755
4	retired	0.385069	0.090909
5	technician	0.151755	0.253731
6	admin.	0.174436	0.245810
7	management	0.210051	0.170604
8	housemaid	0.102894	0.277778
9	entrepreneur	0.116505	0.322581
10	student	0.591837	0.030769
11	unknown	0.200000	0.428571

Bias analysis results for Age groups

	age_group	False Positive Rate	False Negative Rate
0	30-40	0.158148	0.218182
1	40-50	0.130719	0.252560
2	60+	0.661376	0.163399
3	<30	0.218150	0.204473
4	50-60	0.141834	0.215827

[]:

[]: